LOW MEMORY GRAPH TRAVERSAL

by

Samson Bassett

B.Sc. Hon., Thompson Rivers University, 2007.

THESIS SUBMITTED IN PARTIAL FULFILLMENT

OF THE REQUIREMENTS FOR THE DEGREE OF

MASTER OF SCIENCE

IN THE DEPARTMENT

OF

MATHEMATICS

© Samson Bassett 2010

SIMON FRASER UNIVERSITY

July 2010

APPROVAL

Name: Samson Bassett

Degree: Master of Science

Title of thesis: Low memory graph traversal

Examining Committee: Dr. Ralf Wittenberg (Chair)

_____

Dr. Ladislav Stacho

_____

Dr. Petr Lisonek

_____

Dr. Binay Bhattacharya

Date Approved: August 23rd, 2010

Abstract

We provide two traversal algorithms. In the second chapter, we demonstrate that there exists a local orientation for any anonymous, P_3-free graph G with $\delta(G) \geq 2$, such that we can periodically traverse the graph with period $\pi(n) \leq 2n - 2$; this matches the lower bound for the period for general graphs. In the third chapter, we provide a traversal for labelled graphs given that they satisfy two properties, which we define. The properties that we define and use are inspired by the properties of geometric planar graphs.

Dedicated to Rikki.

Acknowledgments

To Mr. Hanson, who taught me how to learn.

*To Dr. Rick Brewster, who showed me the
computational side of mathematics.*

*Finally, to Dr. Ladislav Stacho, whose endless patience,
dedication, and support made this possible.*

Contents

List of Figures

Chapter 1

Introduction

1.1 Motivation and Background

In this thesis we consider the *graph traversal* problem, which asks for a route on a graph that visits every vertex, after starting from an arbitrary vertex. We are interested in finding a graph traversal algorithm, which starts at any given initial vertex, and moves between adjacent vertices, to eventually visit every vertex of the graph at least once. This problem is similar to than the problem of *route discovery*, or *st-connectivity*, which wants to find a route between two given vertices s and t. Clearly, a solution to the graph traversal problem will guarantee a route between any two vertices, and hence a graph traversal can also be used to solve *st*-connectivity. Note that in general, algorithms for these problems will not find the shortest possible paths. Some variations of the graph traversal include additionally requiring the route to end on the same vertex it started on, or visit every edge of the graph.

The model of computation used for a graph traversal is typically some type of mobile agent, or robot, that only knows local information of the graph, (i.e., it cannot see the entire graph). These robots are usually defined as finite automata. The robot visits exactly one vertex at a time, called the current vertex of the robot, and may put a pebble on the vertex (if it has pebbles). The running time (sometimes called time complexity) of a traversal algorithm is an upper bound on the number of edges the robot visits during the traversal. Hence the running time of any traversal algorithm has the trivial lower bound of $\Omega(n)$ for an n vertex graph. The memory used by the

robot, (sometimes called space complexity), which is measured in bits, includes the states, as well as any other memory (such as a bit representation of pebbles) used by the particular robot. For example, a finite automaton with q states uses at least $\log q$ bits of memory, (since it is in one state at a time, and each state has a label of $\log q$ bits). In the following, we further discuss the memory of the robot, depending on the context.

Graph traversal problems come in two categories, based on whether or not the robot can distinguish between vertices of the graph. We say the graph is a *labelled* graph, if each vertex has a unique label, (that is readable by the robot). For example, we can use labels $1, 2, \ldots, n$ for a graph with n vertices. Hence, for labelled graphs, there is a trivial memory lower bound of $\Omega(\log n)$ bits, since $\log n$ bits are required to represent an integer no larger than n, (and since the number of states in the finite automaton is constant). If every vertex of the graph does not have a unique label, we say the graph is *anonymous*. In anonymous graphs, we allow edges at an incident vertex to be distinguishable from one another (by the robot); this is called a *local orientation* of the graph. This is represented by port numbers at each vertex, which are integers between 1 and the degree of the vertex. Since port numbers depend on n, (such as a vertex of degree $n - 1$), the memory is still bounded below by $\Omega(\log n)$.

1.2 Definitions and Notation

A finite simple graph G is an ordered pair $G = (V, E)$, where V is a finite set of vertices, and E is a finite set of edges such that each edge is a two element subset of V. We refer to finite simple graphs as just graphs. For an edge $e = \{u, v\}$ we simply write $e = uv$; we say that e is incident to both u and v, and u is adjacent to v. The *degree* of a vertex v, denoted $d(v)$, is the number of edges incident to v. The minimum degree of a graph G is denoted $\delta(G)$; similarly the maximum degree of G is denoted $\Delta(G)$. A graph is k-regular if every vertex has degree k. A 2-regular graph is a collection of disjoint cycles, and a 3-regular graph is sometimes called a *cubic* graph.

A *uv-walk* in a graph G is an ordered sequence of vertices $u = v_0, v_1, \ldots, v_{k-1}, v_k = v$, such that $v_i v_{i+1} \in E$ for all $0 \leq i \leq k - 1$; these are referred to as edges of the walk. The *length* of this walk is k. A graph is *connected* if there is a uv-walk for every pair of vertices u and v. A graph that is not connected is *disconnected*. The vertices u and v are the *terminal* vertices of the uv-walk. If u and v are understood, we simply refer to this uv-walk as a walk. Non-terminal vertices are called *internal* vertices of the walk. A *path* is a walk that does not allow the repetition of vertices. If P is a uv-path, and s and t are vertices of P, then the st-path containing only edges of P is an *st-subpath* of P, or simply a subpath of P. We say that two distinct paths are *internally disjoint* if they have no internal vertices in common. A *closed walk* C of G is a uv-walk such that $u = v$. A *cycle* is a closed walk that does not allow the repetition of vertices. A cycle in G that contains every vertex of G exactly once is called a *Hamiltonian cycle*. A *Hamiltonian walk* is a closed walk in G that contains every vertex of G at least once.

The *distance* between two vertices u and v is the length of a shortest uv-path. The *diameter* of a graph is the maximum distance between any pair of vertices in the graph. The *square* of a graph G, denoted G^2, is constructed by starting from G, and then adding an edge between each pair of vertices at distance two.

A *subgraph* $G' = (V', E')$ of the graph $G = (V, E)$ is a graph such that $V' \subset V$ and $E' \subset E$. A subgraph F of G is an *induced* subgraph if, for any pair of vertices x, y in F, xy is an edge of F if and only if xy is an edge of G. A graph G that does not

contain F as an induced subgraph is *F-free*. If F is a path of length k, we say that G is P_k-free.

A *tree* is a connected graph that contains no cycles. A *leaf* of a tree is a vertex with degree one. A *rooted tree* is a tree with one vertex r designated as the root of the tree. The *height* of a rooted tree is the length of the longest uv-path, such that u is the root of the tree and v is a leaf of the tree. For any pair of adjacent vertices u and v, we say that u is a *parent* of v the distance from u to r is smaller than the distance from v to r. In this case, we also say that v is a *child* of u. A *spanning tree* of a graph G is a subgraph of G that is a tree, and contains every vertex of G. It is clear that every connected graph has a spanning tree.

A *colouring* of a graph G, is an assignment of colours to the vertices of the graph. The colours are often represented by integers. A colouring is *proper* if adjacent vertices have different colours.

A *vertex cut* of a graph G is a set of vertices whose removal renders the resulting graph disconnected. If a vertex cut contains only a single vertex v, we call v a *cut vertex* of G. Similarly, we say that e is a *cut edge* of G if removing e renders the resulting graph disconnected. We say that the *connectivity* of G is k, or G is *k-connected*, if a minimal vertex cut of G contains k vertices. Hence, G is k-connected if and only if for every pair of vertices u and v in G, there exist k internally disojoint uv-paths.

The *symmetric orientation* of G, denoted $\vec{G}$, replaces each edge uv of G with the two *arcs* (u, v) and (v, u). We say that G is the *underlying* graph of $\vec{G}$. For the purpose of this thesis, we consider the degree of a vertex in $\vec{G}$ to be the degree in the underlying graph. For the arc (u, v), we say that u is the tail (vertex) of the arc, and v is the head (vertex) of the arc. We say that (u, v) is an arc from u to v. If $e = (u, v)$, the *reverse* of e is $e^{-1} = (v, u)$. A *directed uv-path* P is an ordered sequence of vertices $u = v_0, v_1, \ldots, v_{k-1}, v_k = v$, without repetition, such that (v_i, v_{i+1}) is an arc for all $0 \leq i \leq k - 1$. The *reverse* of the directed path P, denoted P^{-1}, is the directed path $v_k, v_{k-1}, \ldots, v_1, v_0$.

If S is a set, the size of S is denoted $|S|$. We denote *total order* on the elements of S by $(S, \leq)$, which is a relation that must satisfy the following properties. If $a, b \in S$, then either $a \leq b$ or $b \leq a$. Furthermore if $a \leq b$ and $b \leq a$, then $a = b$. If $a, b, c \in S$,

and both $a \leq b$ and $b \leq c$, it follows that $a \leq c$. We say that $a < b$ if $a \leq b$ and $a \neq b$.

An *alphabet* is a set of elements called characters. A *string* over the alphabet Σ is a finite ordered collection of characters from Σ.

A *finite automaton* A is a five-tuple $A = (\Sigma, Q, q_0, T, F)$, where:

- Σ is the input alphabet

- Q is a finite set of states

- $q_0 \in Q$ is the initial state

- $T : Q \times \Sigma \to Q$ is the transition function

- $F \subseteq Q$ is the set of accepting states

The execution of a finite state machine is as follows. Let A be a finite automaton Let x be a string of n characters from Σ, i.e., $x = s_1 s_2 \cdots s_n$, where each $s_i \in \Sigma$. Initially, A is in state q_0 and it scans s_1. Then A computes the transition function $T(q_0, s_1) = q_i$ and transitions to state q_i. This repeats until the end of the string is reached, at which point A computes $T(q_j, s_n) = q_k$. If $q_k \in F$, then we say A *accepts* the string x; otherwise A *rejects* x.

The running time, (or time complexity), of an algorithm is an upper bound on the completion time of the algorithm; usually expressed in terms of a number of primitive computations. The memory, (or space complexity), of an algorithm is an upper bound on the amount of storage used by the algorithm, and is often measured in bits. These definitions also apply to finite automata. Hence, the running time of a finite automaton is the number of times the transition function is applied, and the memory is $\log |Q|$.

The following notation, called *big-Oh* notation, is used for comparing functions. If $f(n)$ and $g(n)$ are functions (of real numbers), and there exist constants n_0 and c such that $\forall n > n_0$, $f(n) \leq cg(n)$, then we say that $f(n) = O(g(n))$. For example, $255n^3 + n + 42 + 7 \log n = O(n^3)$. Similarly, we have notation for lower bounds. If there exists constants n_0 and c such that $\forall n > n_0$, $f(n) \geq cg(n)$, then we say that $f(n) = \Omega(g(n))$.

1.3 Related Results

1.3.1 Anonymous Graphs

In this section, we only consider anonymous graphs. Recall that the model of computation is a robot that visits one vertex at a time, and is often modelled as a finite automaton with a constant number of states. The following theorem implies that the graph traversal problem in this context is unsolvable.

Theorem 1.3.1. [16] *For any finite automaton with K states, and any $d \geq 3$, there exists an anonymous planar graph G of maximum degree d with at most $K+1$ vertices that the finite automaton cannot traverse.*

We give a sketch of the proof of Theorem 1.3.1, as shown in [16]. First, take an arbitrary robot with K states. When we say the robot *fails to traverse G*, we mean that for any $t \geq 0$, there exists a vertex v of G such that after the robot has visited t vertices, it has not visited v. A *trap* is a pair (G, u) such that u is a vertex of G, and the robot fails to traverse G when started on u.

The proof will start with an infinite d-regular tree, called T_d, and then start the robot on some vertex u_0 of T_d. By colouring the edges of T_d, and determining how the robot explores the tree, a (G, u_0) trap can be constructed for some graph G obtained from T_d. The graph G may be a subgraph of T_d, but we also allow vertices of T_d to be merged while constructing G. The resulting graph G may have loops and multiple edges. The construction of G is based on the fact that there are only finitely many states, which implies there will always be two vertices, say u and u', that are visited while the robot is in the same state S. Since the vertices have the same degree, and they are unlabelled, and since the robot is in the same state at each vertex, these situations are indistinguishable from one another. We deduce that the robot will periodically return to the state S. Furthermore the path taken between u and u' has some sequence of edge colours L. There are several cases, based on whether or not L is a palindrome, (or is almost a palindrome, in some way). By examining all subcases, and forcing the robot to loop in G in each case, the result follows.

Theorem 1.3.1 implies that graph traversal is impossible if the graph is anonymous, and there is no extra information available to the robot. Next, we wish to find out if

graph traversal is possible if either the robot is allowed to mark the graph, or if the graph is somehow processed prior to the traversal.

Now we suppose the robot has the ability to drop pebbles, (sometimes called mark bits), on the current vertex. The robot can identify whether or not the current vertex has a pebble, and has the option of removing it. If a robot has k pebbles, this adds $O(\log k)$ bits of memory to the robot. The following theorem implies that graph traversal is possible by robot that uses a constant number of pebbles, however the robot requires more than $O(\log n)$ memory.

Theorem 1.3.2. [12] *If G is an anonymous graph with m edges and n vertices, then there exists robot with $k \geq 1$ pebbles that will traverse G while reconstructing it in $O(\frac{mn}{k} - n^2)$ time, and using $O(m \log m)$ memory.*

Theorem 1.3.2 shows that a robot with one pebble is sufficient to traverse any graph. This robot reconstructs the graph while traversing it, and maintains an explored subgraph S and a set of unexplored edges U, which are incident to vertices in S. Hence, the robot knows which edges have been traversed and which edges have not been traversed using U, since it keeps track of the location in both S and G. So at most the robot will have to keep track of m edges in the subgraph S, which means it uses $O(m \log m + k \log k) = O(m \log m)$ since k is constant.

Since the previous result does not use $O(\log n)$ memory, we can ask if there is a lower bound on the memory requirements for graph traversal. This is answered in the following theorem. Note that the lower bound given in Theorem 1.3.3 applies to all graphs.

Theorem 1.3.3. [16] *If G is an anonymous planar graph with diameter D and maximum degree d, a finite automaton requires at least $\Omega(D \log d)$ memory to traverse G.*

Next we examine ways of adding more information to each vertex of the graph, such that it still remains an anonymous graph. One method is to assign colours to the vertices. Theorem 1.3.4 shows that using constant number of colours is sufficient to perform a graph traversal in linear time.

Theorem 1.3.4. [9] *There exists a finite automata with less than 25 states such that for any anonymous graph G with m edges, it is possible to colour the vertices of G with three colours such that the robot will traverse the entire graph G while visiting at most 20m edges.*

It is important to mention that the robot from Theorem 1.3.4 is not allowed to modify the colours of vertices during the traversal. The authors of [9] colour the graph as follows, where the colours can be represented by the integers $0, 1, 2$, (and hence each colour can use two bits per label). First, an arbitrary vertex r is chosen, called the root of the colouring. Then, for each vertex v at distance d from r, v is assigned the colour $d \mod 3$. Notice that this colouring may not be proper, i.e., adjacent vertices are permitted to have the same colour. At the current vertex v, one important task of the robot is to determine how to navigate this tree based on the colouring of the graph. To this end, the robot will examine each edge $e = vu$ incident with v, and determine whether or not u is a parent of v, a child of v, or neither, (by leaving v and then possibly backtracking). With this information, the robot essentially performs a depth-first search of the tree, (for more information on the depth-first search, refer to Section 1.3.2).

The robot uses a constant number of states, however one may wonder whether this 25 state bound could be reduced. To this end, the authors of [9] have another result that shows a robot with only a single state may fail, even if the number of colours used is not constant, (but still less than n).

Theorem 1.3.5. [9] *For any $d > 4$, and any robot with 1 state, there exists a graph G of maximum degree d, such that if G is assigned a colouring with at most $\lfloor \log d \rfloor - 2$ colours, then G cannot be traversed by the robot.*

Recall that anonymous graphs are assumed to have a local orientation, i.e., at each vertex the incident edges have assigned port numbers. In the results above, it is assumed the local orientation has been assigned arbitrarily. This provides the motivation to find out whether it is possible to specifically arrange the port numbers at each vertex to aid the robot in its traversal. The following theorem demonstrates that this is indeed possible.

Theorem 1.3.6. [11] *If G is an anonymous graph on n vertices, there exists a local orientation and a corresponding robot with 1 state that will traverse G in $O(n)$ time. Furthermore, the robot will use at most $10n$ edges during this traversal.*

The basic idea of the proof of Theorem 1.3.6 given in [11] is to examine cycles in G. By constructing subgraphs that have spanning cycles, the authors obtain several cycles, that together visit every vertex of G. The remaining step is to merge the cycles into a single closed walk containing every vertex of G, and they show that this walk contains at most $10n$ edges. The local orientation is arranged so that the robot can simply follow increasing port numbers to stay on this walk. The robot leaves the initial vertex on the edge with the smallest port number. After entering a vertex v at port i, then the robot will leave v on the edge with port $i + 1$. Hence the robot will traverse G as an anonymous graph in linear time with $O(\log n)$ memory, (and without using any pebbles). Since the lower bound for the running time of a graph traversal is $O(n)$, (since each vertex must be visited), the authors of [11] ask whether there is a robot and a local orientation traversal that uses less than $10n$ edges. Theorem 1.3.7 shows this improvement, in the context of a new variant of the graph traversal problem.

The previous result can be seen as a periodic graph traversal, which requires that an algorithm visits every vertex infinitely many times in a periodic manner, where the period is the maximum number of edge traversals performed between two consecutive visits of a generic vertex, denoted by $\pi(n)$. In the context of periodic traversal, another way to state Theorem 1.3.6 is to simply say that $\pi(n) \leq 10n$.

We can observe a trivial lower bound of $\pi(n)$ by considering a tree, where traversal is started at the root. Every edge is visited on a path between the leaf vertex and the root, and since trees are acyclic, this path must be used in both directions by the robot. Since a tree with n vertices has $n - 1$ edges, we have $\pi(n) \geq 2n - 2$. This provides some motivation to assign a local orientation based on a spanning tree of a graph. Indeed, the proof of Theorem 1.3.7 uses a spanning tree, resulting in a shorter period than Theorem 1.3.6.

Theorem 1.3.7. [21] *If G is an anonymous graph on n vertices, there is a local orientation and a robot with 3 states that will periodically traverse G with period*

$\pi(n) \leq 4n - 2$.

The authors of [21] show this by starting with an arbitrary spanning tree T of G with root r. For each vertex $v \neq r$ the smallest ports will be assigned to edges in T, then all other ports will be on edges not in T. Port 1 will always be on the edge of the unique path from v to r in T. A proof by induction is used to show $\pi(n) \leq 4n - 2$, by removing an edge of T that has port 1 at one of its endpoints. This result clearly uses $O(\log n)$ memory, however it is interesting to notice that the robot in this result uses 3 states, as opposed to the robot from Theorem 1.3.6 which uses only 1 state. The author of [21] conjectures that the best possible upper bound of $\pi(n)$ is $4n - c$ for general graphs, where c is a constant. The following result refutes this conjecture.

Theorem 1.3.8. [17] *If G is an anonymous graph on n vertices, there is a local orientation of G and a robot with 11 states that will periodically traverse G, with period $\pi(n) \leq 3.75n - 2$.*

We give sketch of the proof of Theorem 1.3.8 from [17]. The main idea is to expand on the work of [21] by further examining the spanning tree T of G. The goal of this paper is to provide a local orientation that allows the robot to recognize vertices and subtrees with specific properties, and move accordingly. They introduce the concept of an *extended leaf*, which is an edge uv of T such that v is a leaf, and u is not incident to any leaves other than v. Similarly a *paired extended leaf* is an extended leaf connected to a vertex whose children contain at least two extended leaves. Next they arrange the port numbers at each vertex, such that after the robot visits a leaf or an extended leaf, it will always visit a paired extended leaf. Furthermore, the robot will traverse some edges of G that are not in T. We say that a penalty is paid at v if the robot leaves v on an edge that is not in T. To prove correctness, and the period, the authors colour certain vertices based on the above leaf concepts, and based on a certain subtree of T called the backbone. By examining this backbone, and the vertex colours, the penalties paid (including both directions of a given edge) is shown to be at most $\frac{7}{4}n$. Also, at most $2n - 2$ edges of the tree will be used during the traversal. Hence the bound $\pi(n) \leq 3.75n - 2$ is obtained. Notice that the number of states of the robot has again increased by a constant factor (from 3 to 11), between Theorem 1.3.7 and Theorem 1.3.8. These different states of the robot from Theorem

1.3.8 are used to perform several different subroutines, such as searching for the root, backtracking, doing something specific in regards to extended leaves, and so on. It seems natural to wonder whether or not this trend would continue, i.e., what is the tradeoff between the number of states of the robot and the period of the traversal performed by the robot.

The authors of [17] note that the idea of using subtrees has provided strong results, as opposed to the cycle based proof of Theorem 1.3.6. They ask whether the period can be further improved by using spanning trees, or using a different technique. The problem of finding the minimal upper bound on $\pi(n)$ remains open.

1.3.2 Labelled Graphs

In this section, we only consider labelled graphs. In this context, we often say that the algorithm traverses an edge or visits a vertex. Two well known solutions to the graph traversal problem for labelled graphs are the Breadth-First Search (BFS) and the Depth-First Search (DFS) algorithms. We present these algorithms using the queue and stack data structures, respectively, (so the algorithm will be storing vertex labels in these data structures).

BFS(Graph G, vertex s)

  1: Set $S = \{s\}$
  2: Queue $Q = $ all neighbors of s in any order
  3: **while** Q is not empty **do**
  4:     dequeue vertex v from the front of Q
  5:     **if** v is not in S **then**
  6:       add v to S
  7:       enqueue neighbors of v onto the end of Q in any order
  8:     **end if**
  9: **end while**

The BFS algorithm will move out from the starting vertex to every neighbour at once. Then from each of those new vertices, every unvisited neighbour is visited. So first all vertices at distance one from the start are visited, then all vertices at distance two, and so on.

The DFS algorithm works differently; it first moves as far as possible away from the start untill there are no more adjacent vertices that have not been visited. Then it will back up one vertex at a time, and attempt to find another route with unvisited vertices, and perform the same action.

DFS(Graph G, vertex s)

1: Set $S = \{s\}$
2: Stack T = all neighbors of s in any order
3: **while** T is not empty **do**
4:     pop vertex v from the end of T
5:     **if** v is not in S **then**
6:         add v to S
7:         push neighbors of v onto the end of T in any order
8:     **end if**
9: **end while**

Both the BFS and the DFS algorithms use $O(n \log n)$ memory, since in the worst case every vertex is saved in the data structure, (each of which has a label that uses $\log n$ bits). The running time of both the BFS and DFS algorithms is $O(m + n)$. To see this, consider a tree, so in the worst case, every vertex and every edge is used by these algorithms. Notice that neither the running time nor the memory used are optimal for these two algorithms. Recall that the lower bounds for graph traversal are $\Omega(n)$ time and $\Omega(\log n)$ bits of memory. Most soultions focus on either minimizing the running time at the expense of memory, or minimizing the memory at the expense of time. We focus on the latter. First, we mention the following result showing that there is a tradeoff between memory usage and running time.

Theorem 1.3.9. [14] *Any n vertex graph can be traversed in memory S and time T, such that $ST \leq mn(\log n)^K$, where K is a constant.*

So, Theorem 1.3.9 implies there is a spectrum of compromises between memory and running time for the graph traversal problem. Next, we investigate whether graph traversal is possible using at most $O(\log n)$ memory. For example, a constant number of finite automata, each with a constant number of states (and no external memory) together use at most $O(\log n)$ memory. However, Theorem 1.3.10 shows that such a team is not sufficient to traverse all graphs.

Theorem 1.3.10. [23] *A finite team of finite automata (each with a constant number of states) cannot traverse all planar cubic graphs.*

Note that the geometry of the planar embedding of the graph in Theorem 1.3.10 is not used by the algorithm, (i.e., it is not a geometric graph). Theorem 1.3.10 implies that once we fix some collection of n finite automata, there exists a graph (which may have more than n vertices) that this team of automata cannot traverse. This seems to imply that knowing the size of the graph in advance may be useful, and indeed some results in the following assume this.

Next, we consider a model of computation, introduced in [7], that is similar to a team of finite automata. A *Jumping Automaton for Graphs*, or JAG, is a finite state automaton that is equipped with a set of pebbles that are all initially placed on some vertex. The JAG does not traverse the graph itself, but it can move one pebble to an adjacent vertex, and teleport one pebble to the location of any other pebble. A JAG with P pebbles and Q states, (traversing a graph with n vertices), uses $P \log n + \log Q$ bits of memory [7]. Notice that if a JAG is restricted to a constant number of pebbles and a constant number of states, it uses $O(\log n)$ memory. If a JAG has n pebbles, it can perform a BFS or a DFS of any graph with n or fewer vertices. Hence, a JAG can traverse a graph if it knows an upper bound on the number of vertices the graph can have. However, if we suppose the JAG knows the size of the graph in advance, the following result shows that the JAG will traverse the graph using less memory than a BFS or DFS, since $O((\log n)(\log n) + \log(n^4)) = O(\log^2 n)$.

Theorem 1.3.11. *For any n, there is a JAG with $O(\log n)$ pebbles and $O(n^4)$ states that can traverse all graphs with n vertices in polynomial time. This JAG uses $O(\log^2 n)$ memory.*

Furthermore, the following result implies there is a tradeoff between memory and running time, for JAGs that know n.

Theorem 1.3.12. [13] *For every $z \geq 2$, and any fixed n, a JAG with at most $\frac{\log n}{28z \log \log n}$ pebbles and at most $2^{\log^z n}$ states requires at least $\Omega(n2^{\frac{\log n}{\log \log n}})$ time to traverse a cubic graph on n vertices.*

For more results that investigate this tradeoff, see [3] and [5]. The authors of [3] conjecture that for any JAG that traverses a graph using S memory and T time, $ST \geq Kmn$, where K is a constant, (refer to Theorem 1.3.9 for an upper bound).

The previous results show that if a JAG has some upper bound on the number of vertices of a given graph, it can traverse the graph. However, the following result implies that there is no fixed JAG that can traverse all graphs (with an unbounded number of vertices).

Theorem 1.3.13. [7] *For every JAG with a fixed number of pebbles and states, there is a graph with maximum degree 3 that cannot be traversed by the JAG.*

Theorem 1.3.10 and Theorem 1.3.13 seem to imply that, for general labelled graphs, traversal may not be possible if we are restricted to only $O(\log n)$ bits of memory. So we can ask what is the minimum number of memory bits required for a robot to traverse any labelled graph. The following results are related to this question.

Theorem 1.3.14. [2] *Any labelled graph on n vertices can be traversed in polynomial time and $O(\log^2 n)$ memory.*

The memory bound in Theorem 1.3.14 matches the one given by Theorem 1.3.11 using the JAG model. The following result improves this bound.

Theorem 1.3.15. [22] *Any labelled graph on n vertices can be traversed in polynomial time and $O(\log^{1.5} n)$ memory.*

There is another slight improvement of the memory bound in the following theorem, but it uses non-polynomial time in the worst case.

Theorem 1.3.16. [1] *Any labelled graph on n vertices can be traversed in non-polynomial time and $O(\log^{4/3} n)$ memory.*

Next, we want to see if we can make an assumption about the graph, such that a robot can traverse the graph using only $O(\log n)$ bits of memory. To this end, we consider the following embedded graphs. A *geometric graph* is an embedded graph, where each vertex is associated with the coordinates of a unique point in Euclidean space. Furthermore, if the robot is visiting some vertex, it can read the coordinates of that vertex in the embedding. An embedded planar graph, where each vertex has cartesian coordinates from the plane, is an example of a geometric graph. It turns out this assumption is sufficient to traverse the graph using only $O(\log n)$ memory.

Theorem 1.3.17. [10] *Let G be a geometric planar graph with n vertices. Then G can be traversed in $O(n^2)$ time and $O(\log n)$ memory.*

The proof of 1.3.17 from [10] simulates a DFS of a spanning tree of the dual graph of G. Each vertex of the dual graph represents a face of the embedded graph. Each face will be traversed in some cyclic order, and at certain edges the algorithm will start traversing a different face that shares that edge. Each edge is assigned a unique rational number based on the geometry, which implies each face has a minimum edge. These minimum edges are used to enter and leave the faces of the embedding, which implies that the faces are adjacent in the spanning tree. The following result extends Theorem 1.3.17, giving an improved running time.

Theorem 1.3.18. [4] *Let G be a geometric planar graph with n vertices. Then G can be traversed in $O(n \log n)$ time and $O(\log n)$ memory.*

The proof of Theorem 1.3.18, given in [4], uses a similar technique as the authors of [10]. Now that we have seen some results for planar graphs, we investigate non-planar geometric graphs. The authors of [6] extended Theorem 1.3.18 to a larger class of graphs, called quasi-planar graphs; see Theorem 1.3.19. A *quasi-planar* graph, is an embedded graph G with an induced planar subgraph P, such that the outer face

of P in the embedding does not contain any vertex or edge of $G - P$, and no edge of P is crossed by any edge of G. Any graph that has at least one cycle has the following quasi-planar embedding: simply embed the cycle as the outer face, and place every other vertex on the interior of the cycle.

Theorem 1.3.19. [6] *Let G be a graph with n vertices that has a quasi-planar embedding that satisfies the Left-Neighbor Rule. Then G can be traversed in $O(m \log m)$ time and $O(\log n)$ memory.*

The Left-Neighbour Rule stated in Theorem 1.3.19 requres that for any vertex v that is not in the planar subgraph of the embedding, there is a vertex u whose x-coordinate is smaller than the x-coordinate of v. Notice that Theorem 1.3.19 traverses the graph using a specific embedding of the graph, i.e., the embedding cannot be given arbitrarily. The authors of [6] prove this theorem using similar techniques to the ones shown in [4] and [10]. The question of whether we can traverse all geometric graphs using $O(\log n)$ memory, remains open.

1.4 Organization of this Thesis

In the second chapter, we demonstrate a periodic traversal of any P_3-free graph G with $\delta(G) \geq 2$. We achieve an upper bound on the period of $\pi(n) \leq 2n - 2$, which matches the lower bound, and is less than the period given in Theorem 1.3.8. One of the techniques we use is based on a Hamiltonian cycle of the square of the graph.

In the third chapter, our goal is to demonstrate a graph traversal similar to the one from Theorem 1.3.19 that does not require the specific two dimensional quasi-planar embedding. Our traversal also runs in $O(m \log m)$ time and uses $O(\log n)$ memory. This is done by generalizing the properties that are used by the authors of [6] in the proof of Theorem 1.3.19, and these properties are stated independently of any embedding.

Chapter 2

Periodic Traversal

In this chapter, we demonstrate how to periodically traverse any graph G that is P_3-free such that $\delta(G) \geq 2$. This is accomplished by assigning port numbers at each vertex, and then starting a robot at any initial vertex. There are two types of graphs for which we assign port numbers differently, after which the same robot will traverse the graph periodically.

Our first lemma identifies these two cases, by determining what happens when a P_3-free graph is not 2-connected.

Lemma 2.0.1. *Let G be a graph with n vertices such that G is P_3-free. Then either G is 2-connected, or $\Delta(G) = n - 1$.*

Proof. Suppose that G is not 2-connected, which implies G has a cut vertex v. Suppose to the contrary that $\Delta(G) < n - 1$, so $d(v) < n - 1$. Hence there is a vertex x that is not adjacent to v. Let G_x be the component of $G - v$ that contains x. Since v is a cut vertex of G, there is a vertex y in component G_y of $G - v$ such that y is adjacent to v and $G_y \neq G_x$. Since G_x is connected, we may assume that x is at distance 2 from v, so there is a vertex x' that is adjacent to both v and x.

Let P be the path that visits the vertices y, v, x', and x. Since $G_y \neq G_x$, y is not adjacent to either x or x'. Hence P is an induced path of G with three edges, since by supposition x and v are not adjacent. This is a contradiction, since G is P_3-free. Therefore $\Delta(G) = n - 1$. □

We examine G in two cases in Section 2.2 and Section 2.3, depending on whether

or not G has a cut vertex. First, we define the robot that will perform the periodic traversal in the next section.

2.1   The Robot

In this section, we will define our robot to be a Mealy automata, similar to the proofs in [17]. The robot is a type of finite automata called a Mealy automata for graphs, which is a five-tuple $M = (\Sigma, S, \alpha, v, T)$, where $\Sigma = [1, d(v)] \times V$ is the alphabet, $S = \{\alpha\}$ is the set of states, α is the initial state, v is the initial vertex, and T is the transition function. The input and output of the transition function T consists of one element of the alphabet, and one state; it is defined below.

The robot will move between adjacent vertices along arcs, according to the transition function. To *leave* the current vertex v on port number p, the robot will traverse the arc $e = (v, u)$ such that e has the port number p. The robot then moves to the vertex u, and this vertex becomes the current vertex of the next step. We also say that the robot *entered* u on the port q, where q is the port number of $e^{-1} = (u, v)$.

The robot will always be in a given configuration (p_1, v_1, α), where p_1 is the entering port, v_1 is the current vertex, and α is the current state. It can examine all the arcs incident to v_1, and hence all the ports seen at v_1. Using this information, the robot applies the transition function to decide which port it will leave the current vertex on, and which state to change to. This ultimately also determines the next current vertex. The transition function is of the form $T((p_1, v_1), \alpha) = ((p_2, v_2), \alpha)$, where p_1 is the entering port, v_1 is the current vertex, p_2 is the leaving port, v_2 is the new current vertex such that v_2 is the head of the arc incident to v_1 with port p_2. For simplicity, we will write $T(p_1, v_1, \alpha) = (p_2, v_2, \alpha)$ instead of $T((p_1, v_1), \alpha) = ((p_2, v_2), \alpha)$. We say that (p_2, v_2, α) is the new current configuration of the robot. Interestingly, this particular robot always stays in the initial state α. However, for formality, we include α in the configuration and the transition function. The transition function T is defined as follows.

$$
\begin{aligned}
T(p, v, \alpha) &= (q, u, \alpha) : p < d(v), (v, u) \text{ has port } p + 1, \text{ and } (u, v) \text{ has port } q \\
T(d(v), v, \alpha) &= (q, u, \alpha) : (v, u) \text{ has port } d(v) \text{ and } (u, v) \text{ has port } q
\end{aligned}
$$

Each application of the transition function changes the current configuration of

the robot. When the robot enters a vertex v on port $p < d(v)$, the transition function instructs the robot to leave v on the arc (v, u) that has port $p + 1$. If v was entered on port $d(v)$, the robot will leave on the reverse of the entering edge, which has port $d(v)$.

2.2 The graph has a cut vertex

For the remainder of this section, we assume G is a P_3-free graph on n vertices such that $\delta(G) \geq 2$, and G is not 2-connected. Hence by Lemma 2.0.1, $\Delta(G) = n - 1$.

First, we fix three specific vertices that will be assigned port numbers in a more specific way than the rest of the vertices. Let v^* be a vertex of G such that $d(v^*) = n - 1$, and let $x^* \neq v^*$ be any other vertex. Since $\delta(G) \geq 2$, $d(x^*) \geq 2$, so there is a vertex $y^* \neq v^*$ adjacent to x^*. These are called the star vertices of G. Since $d(v^*) = n - 1$, these three vertices v^*, x^*, and y^* form a cycle in G. This is called the star-cycle of G.

For each vertex v of G, we assign port numbers to each edge incident to v. Recall that in a local orientation of G, each edge e has two port numbers, one at each incident vertex. If $e = uv$, we say that port p is assigned to uv if that p is assigned to e at u; similarly we say that port p is assigned to vu if p is assigned to e at v. For our robot, we need to avoid the situation where uv has port $d(u)$, and vu has port $d(v)$. This is accomplished by executing the Star-Numbering procedure once on every vertex of G.

---

procedure Star-Numbering(v)

1: **if** $v = v^*$ **then**
2:     Assign $d(v)$ to vx^*
3:     Assign 1 to vy^*
4: **end if**
5: **if** $v = x^*$ **then**
6:     Assign $d(v) - 1$ to vv^*
7:     Assign $d(v)$ to vy^*
8: **end if**
9: **if** $v = y^*$ **then**
10:     Assign $d(v) - 1$ to vx^*
11:     Assign $d(v)$ to vv^*
12: **end if**
13: **if** $v \neq v^*$ and $v \neq x^*$ and $v \neq y^*$ **then**

14:     Assign $d(v)$ to vv^*
15: **end if**
16: arbitrary-ports(v)

Clearly, once the Star-Numbering procedure has been executed for every vertex, this produces a local orientation of G. The following lemma shows that, with this local orientation, every vertex of G will be visited by the robot.

Lemma 2.2.1. *Let G be a P_3-free graph on n vertices such that $\delta(G) \geq 2$, $\Delta(G) = n-1$, and suppose the robot starts in configuration $(d(v), v, \alpha)$, where v is any vertex of G. Then the robot will visit every vertex of G, while performing at most $2n - 3$ edge traversals between two consecutive visits of v.*

Proof. Let v be any vertex of G. Suppose for simplicity v is not a star vertex of G, (if it is, the proof is similar).

The robot will enter v^* on some port $1 < p < d(v^*)$. Hence the robot will leave on port $p + 1$ and enter the vertex $u \neq v$. Since uv^* has port $d(u)$, the robot will return to v^* back along the same edge, and hence it enters v^* on port $p + 1$.

This process repeats until eventually the robot visits the star-cycle for the first time, traverses it, then eventually leaves v^* on port 1. Eventually, the robot will return to v. There are three edges in the star-cycle, each visited once. Outside of this, there are $n - 3$ edges incident to v^*, each of which are visited twice. This gives the total of $3 + 2(n - 3) = 2n - 3$.

$\square$

This lemma is used in the proof of the Theorem 2.4.1 in Section 2.4.

2.3 The graph is 2-connected

In this section, we assume G is P_3-free and 2-connected. We have some simple definitions based on $\vec{G^2}$, (which is the symmetric orientation of the square of the graph G). An arc e of $\vec{G^2}$ is called a *real* arc if e is an arc of $\vec{G}$, otherwise e is a *virtual* arc. A directed path in $\vec{G^2}$ that contains only real arcs is a *real path*, and similarly a directed path containing only virtual arcs is a *virtual path*.

Suppose that H is a Hamiltonian cycle of G^2 with the minimum number of virtual arcs. Such a cycle exists since G is 2-connected, as demonstrated in the well known result of Fleishner in [15]. For our purposes, we will assume the vertices of H are placed on a circle in the order they appear on H. Moreover, we will assume a counterclockwise orientation of H, which makes H a directed cycle in $\vec{G^2}$. We will abuse notation and consider H as both a subgraph of $\vec{G^2}$, and a directed cycle in $\vec{G^2}$. The following lemma is a simple observation based on the minimality of H.

Lemma 2.3.1. *If (u, v) is any virtual arc of H, then there does not exist a (real or virtual) arc (x, y) of H such that (x, u) and (y, v) are real arcs of $\vec{G^2}$.*

Proof. Suppose there does exist such an arc (x, y) of H, such that both (x, u) and (y, v) are real arcs. Let P be the directed vx-path in H, and let Q be the directed yu-path in H. We can now construct a Hamiltonian cycle H' of $\vec{G^2}$ by following P, then the arc (x, u), then Q^{-1}, and finally the arc (y, v). See Figure 2.1. Since (u, v) is a virtual arc, H' has at least one fewer virtual arcs than H, a contradiction. $\square$
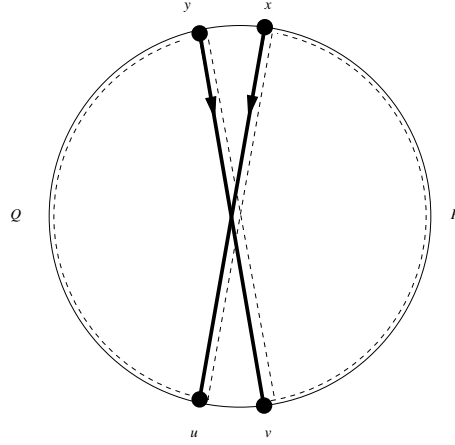
Figure 2.1: *The Hamiltonian cycle H is represented by a circle, and similarly H' by a dashed line. The real arcs are represented by thicker lines.*

2.3.1 Relay Vertices and Wedges

For each virtual arc (u, v) of H, we assign a real directed uv-path of length two. This assigned path is called the *relay path* of the arc (u, v). The fact that a relay path exists for every virtual arc follows from the definition of $\vec{G}^2$. The unique internal vertex of the relay path is called the *relay vertex* of the arc (u, v). Note that more than one such real directed path may exist for a virtual arc. The relay paths can be chosen arbitrarily, it is only important that they are fixed.

A vertex may be the relay vertex for more than one virtual arc. Hence we will group together virtual arcs that are consecutive on H if they have the same relay vertex. If P is a maximal virtual (directed) path in H containing more than one vertex, such that all virtual arcs of P share a common relay vertex w, then the subgraph W of $\vec{G}^2$ consisting of P, and the relay path of every arc of P, is called a *wedge*; see Figure 2.2 for an example. We call P the *virtual path of the wedge W*. Similarly, arcs of P are called *virtual arcs of the wedge W*, and the vertex w is called the *relay vertex of the wedge W*. Real arcs of W are called *ribs* of W. If we suppose the terminal vertices of P are u and v such that P is a uv-path, we refer to u as the *left external vertex* of W, and v as the *right external vertex* of W; both are called *external vertices* of W. Internal vertices of P are called *internal vertices* of the wedge. The rib (u, w) is the *left external rib* of W, and the rib (w, v) is the *right external rib*

of the W; they are both referred to as *external ribs* of W. Real arcs of W that are incident to an internal vertex are called *internal ribs* of W. We define the *size* of a wedge W, denoted $s(W)$, as the number of vertices of its virtual path.
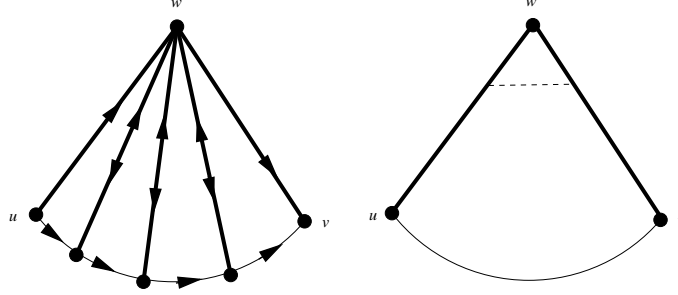
Figure 2.2: *A wedge of size five. The ribs are real arcs of $\vec{G^2}$ and are represented in this figure by thicker lines. The virtual path consists of virtual arcs of $\vec{G^2}$ and are represented by regular lines.*

Note that ribs of a wedge are always real arcs of $\vec{G^2}$. The following lemma characterizes how two wedges can interact.

Lemma 2.3.2. *Suppose W_1 and W_2 are distinct wedges. Then*

(i) The wedges W_1 and W_2 have no virtual arcs in common.

(ii) The wedges W_1 and W_2 have no ribs in common.

(iii) If e is a rib of W_1, e^{-1} is not a rib of W_2.

Proof. Let W_1 and W_2 be wedges with relay vertices w_1, w_2 and virtual paths P_1, P_2, respectively.

To show (i), suppose to the contrary that e is a virtual arc of both W_1 and W_2. Now w_1 is the relay vertex of e, since e is a virtual arc of W_1. But e is also a virtual arc of W_2, so w_2 is the relay vertex of e. Hence $w_1 = w_2$, since each virtual arc has exactly one relay vertex. Now, since we supposed that $W_1 \neq W_2$, it follows that $P_1 \neq P_2$. Since e is an arc of both P_1 and P_2, it must be the case that either one of the paths is a subpath of the other, contradicting the maximality of the smaller path; or P_1 and P_2 overlap, in which case neither path is maximal, again a contradiction. Therefore W_1 and W_2 have no virtual arcs in common.

We prove (ii) in two cases, first suppose $w_1 = w_2$. Suppose to the contrary that e is a rib of both W_1 and W_2. If e is an internal rib of either W_1 or W_2, then v is

incident to a virtual arc that is in both P_1 and P_2, which contradicts (i). So e is an external rib of both W_1 and W_2. Suppose that e is the right external rib of W_1, (the case where e is a left external vertex of W_1 is symmetric). So $e = (w_1, v)$ such that v is the right external vertex of W_1, (and $v \neq w_1$). If e is also the right external rib of W_2, then v is incident to a virtual arc that is in both P_1 and P_2, which contradicts (i). Therefore e is the left external rib of W_2. Let u be the left external vertex of W_2, so $e = (u, w_2)$ is the left external rib of W_2. This is a contradiction, since the head of e is $v = w_2 = w_1$, and $v \neq w_1$.

Now suppose $w_1 \neq w_2$. Suppose to the contrary the real arc e of $\vec{G^2}$ is a rib of both W_1 and W_2. This implies e is incident to both w_1 and w_2. Without loss of generality, let $e = (w_1, w_2)$. First we will show that w_1 is an external vertex of W_2, and similarly w_2 is an external vertex of W_1. Suppose to the contrary that w_1 is an internal vertex of W_2, (we obtain a similar contradiction if w_2 is an internal vertex of W_1). Therefore there are two virtual arcs of W_2 incident with w_1, say (u, w_1) and (w_1, v). There is at least one virtual arc of W_1 incident with w_2, call it h. Suppose that $h = (x, w_2)$. We know that (x, w_1) is a real arc since it is a rib of W_1, similarly (w_2, v) is a real arc since it is a rib of W_2, see Figure 2.3. This contradicts Lemma 2.3.1. If $h = (w_2, x)$ we get a similar contradiction. Therefore w_1 is an external vertex of W_2, and similarly w_2 is an external vertex of W_1.

It follows that e is an external rib of both W_1 and W_2. Since w_1 is the tail of e, e is the right external rib of W_1, and similarly e is also the left external rib of W_2. Therefore (w_1, v) is a virtual arc of the wedge W_2, and (x, w_2) is a virtual arc of the wedge W_1, see Figure 2.3. So (x, w_1) and (w_2, v) are ribs of their respective wedges. As above, since all ribs are real arcs, this contradicts Lemma 2.3.1. This concludes the proof of (ii).

To show (iii), let $e = (u, v)$ be a rib of W_1. Suppose to the contrary that $e^{-1} = (v, u)$ is a rib of W_2. If e is an internal rib of W_1, then e^{-1} is also an internal rib of W_1, which contradicts part (ii). So e is an external rib of W_1; similarly e^{-1} is an external rib of W_2. Suppose without loss of generality that e is the right external rib of W_1, (the case e is the left external rib of W_1 is symmetric). So $u = w_1$, and v is the right external vertex of W_1.

If e^{-1} is the left external rib of W_2, then $u = w_2$ and v is the left external vertex
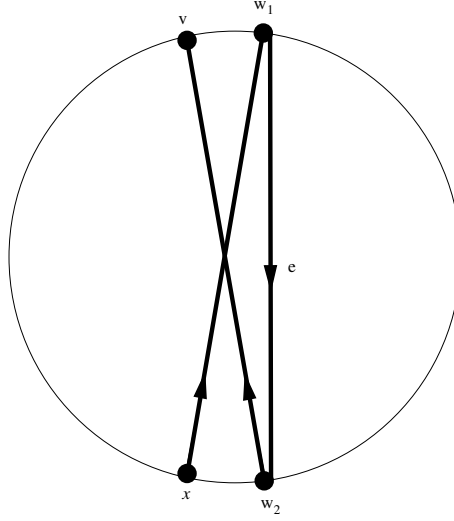
Figure 2.3: *The relay paths of the virtual arcs (x, w_2) and (w_1, v) both contain e.*

of W_1. So the union of P_1 and P_2 is a virtual path such that each (virtual) arc of this new path has relay vertex u, m which contradicts the maximality of P_1 (and P_2).

Next suppose $e^{-1} = (v, u)$ is the right external rib of W_2, and hence $v = w_2$. Hence $e = (w_1, w_2)$, and $e^{-1} = (w_2, w_1)$. Since w_2 is the right external vertex of the wedge W_1, there is a vertex x such that (x, w_2) is a virtual arc of W_1, and (x, w_1) is a rib of W_1. Similarly, there is a vertex y such that (y, w_1) is a virtual arc of W_2, and (y, w_2) is a rib of W_2.

Let P be the path in G containing the edges xw_1, w_1w_2, w_2y. Since all these vertices are distinct, P is a path of length three in G. To show that P is an induced subgraph of G, first notice that xw_2 and yw_1 are clearly virtual edges. If xy is a real edge, then (x, y) and (y, x) are real arcs, which contradicts Lemma 2.3.1. So xy is not a real edge, which implies P is an induced subgraph of G, a contradiction since G is P_3-free. Therefore e^{-1} is not a rib of W_2. $\qquad\square$

We can see that an external rib of a wedge cannot be an arc of H. However, an external rib may be incident to consecutive vertices of H. Indeed, if (w, v) is the right external rib of a wedge W such that its inverse (v, w) is an arc of H, then we say that the wedge W is a *right tied* wedge; similarly if (u, w) is the left external rib of W such that (w, u) is an arc of H, then W is a *left tied* wedge. A wedge that is neither

left tied nor right tied is called a *free wedge*. Next, we show that there are not exist a wedge that is both left tied and right tied.

Lemma 2.3.3. *There does not exist a wedge of $\vec{G}^2$ that is both left tied and right tied.*

Proof. Suppose to the contrary that the wedge W is both left tied and right tied. It is clear that W contains all the vertices of G. Let w be the relay vertex of W, u be the left external vertex of W, and v be the right external vertex of W. If W contains only these three vertices, then (u, v) is a virtual arc of W. Hence there is no arc joining u and v in $\vec{G}$. However, this implies that w is a cut vertex in G, a contradiction since G is 2-connected. Now suppose G has at least four vertices.
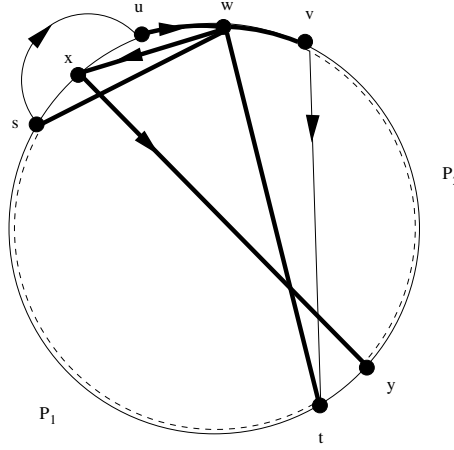

Figure 2.4: *The vertex w is the relay vertex of a wedge that is both left tied and right tied.*

Let P be the virtual path of W, so P is a directed uv-path. Let $x \neq v$ be the vertex adjacent to u on P. Such a vertex exists since $\vec{G}$ has at least four vertices. Since G is 2-connected, $d(x) \geq 2$, so there is at least one real arc other than the ribs (w, x) and (x, w) incident to x, say $e = (x, y)$. (Recall that we calculate the degree of a vertex with respect to G). Since (u, x) is a virtual arc, $u \neq y$, (but it is possible that $y = v$). Since (x, y) is a real arc, there exists at least one internal vertex of W that lies between x and y on the path P; see Figure 2.4. Let s be the vertex following x on P, and t be the vertex preceding y on P, (it is possible that $s = t$). Let P_1 be the st-subpath of P, and P_2 be the yv-subpath of P, (either of which may not contain

any arcs). Since there is a real path of length two connecting s and u, (s, u) is either a real arc or a virtual arc of $\vec{G}^2$. Similarly (v, t) is an arc of $\vec{G}^2$, (that is either real or virtual).

Let H' be the cycle in $\vec{G}^2$ following (u, w), (w, x), (x, y), P_2, (v, t), P_1^{-1}, and (s, u), see Figure 2.4. Clearly, H' is a Hamiltonian cycle of $\vec{G}^2$. This is a contradiction since H contains only two real arcs, (namely (v, w) and (w, u)), while H' contains at least three real arcs, (namely (u, w), (w, x), and (x, y)). Therefore W cannot be both left tied and right tied. $\qquad\square$

For any vertex v, there may be several distinct wedges with relay vertex v; in the following we identify these wedges. If v is the relay vertex of $k \geq 1$ wedges, we order these wedges at v based on their occurrence along H. The *list of ordered wedges* $W_1, W_2, \ldots, W_k$ *at* v is the ordered list of all wedges with the common relay vertex v such that for W_i and W_j with $i < j$, the virtual path of W_i occurs before the virtual path of W_j when following H starting from v; for example see Figure 2.5. When the list of ordered wedges at a vertex is understood, we will write s_i for $s(W_i)$, (recall that $s(W)$ denotes the size of the wedge W). If e is a real arc incident with v, we say that e is an *isolated arc* of v if neither e nor e^{-1} are ribs of any wedge in the list of ordered wedges at v.
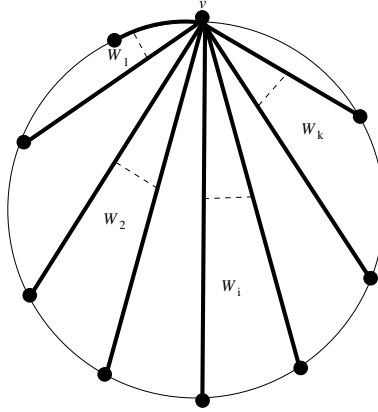
Figure 2.5: *The list of ordered wedges at v is $W_1, W_2, \ldots, W_k$. In this example, W_1 is a left tied wedge.*

In the Section 2.1, we will show how to turn the Hamiltonian cycle H in $\vec{G}^2$ into a Hamiltonian walk H^* in $\vec{G}$. To this end, we specify two arcs for each vertex, called

the incoming and outgoing arcs of the vertex.

Definition 2.3.4. *Given a vertex v, let (x, v) and (v, y) be the two arcs of H incident with v.*

- *The incoming arc of v is either (x, v) if it is real, or the real arc (w, v), where w is the relay vertex of (x, v).*

- *Similarly, the outgoing arc of v is either (v, y) if it is real, or the real arc (v, w), where w is the relay vertex of (v, y).*
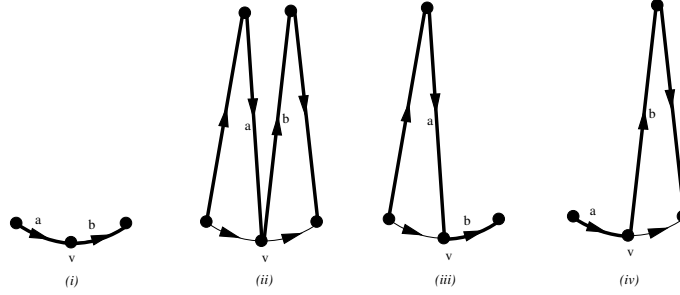
If e is the incoming arc of v and e^{-1} is the outgoing arc of v, we say that e is a backtrack arc of v, (in which case e^{-1} is also a backtrack arc of v).

Note that the incoming and outgoing arcs of a vertex are always real arcs. By definition, each vertex has exactly one incoming arc, and one outgoing arc. Since the relay vertices of all virtual arcs are fixed, the incoming and outgoing arcs are also fixed for every vertex of $\vec{G}$. In general, incoming and outgoing arcs of a vertex may or may not be isolated arcs of that vertex. In the following we let a and b be the incoming and outgoing arcs of the vertex v, respectively

We first consider cases when $a \neq b^{-1}$, i.e., v does not have a backtrack arc; there are four possible configurations. If both a and b are real arcs of H, then they form a real path of length two along H containing v as an internal vertex, see Figure 2.6(i). If v is incident to two virtual arcs of H, then both a and b are arcs of disjoint relay paths, as shown in Figure 2.6(ii). Otherwise, exactly one of a or b is an arc of H, as shown in Figures 2.6(iii) and 2.6(iv). In the following lemma, we investigate how these arcs interact with wedges when they are not isolated arcs of v.

Lemma 2.3.5. *Let a be the incoming arc of v, and b be the outgoing arc of v, such that $a \neq b^{-1}$. Then:*
(i) If a is not an isolated arc of v, then a^{-1} is the right external rib of a right tied wedge W with relay vertex v.
(ii) If b is not an isolated arc of v, then b^{-1} is the left external rib of a left tied wedge W relay vertex v.

Figure 2.6: *The four of the possible configurations where a is the incoming arc of v, b is the outgoing arc of v, and $a \neq b^{-1}$.*

Proof. We prove only (i); the proof for (ii) is symmetric. Suppose a is not an isolated arc of v. It follows that either a or a^{-1} is a rib of some wedge with relay vertex v. First we show that $a = (u, v)$ cannot be a rib of any wedge with relay vertex v.

To this end, suppose to the contrary that a is a rib of the wedge W with relay vertex v. It follows that a is not an arc of H, because wedges cannot contain any real arcs of H. Hence, by Definition 2.3.4, u is the relay vertex of the virtual arc (x, v), (and $u \neq v$). So u is the relay vertex of some wedge W' containing (x, v). But now a is a rib of both W and W'. Clearly $W \neq W'$, since they have different relay vertices. This contradicts Lemma 2.3.2(ii). Therefore a is not a rib of any wedge with relay vertex v.

Since a is not an isolated arc of v, we conclude that there exists a wedge W with relay vertex v such that a^{-1} is a rib of W. Since a is the incoming arc of v, we have two cases by Definition 2.3.4. First, suppose that a is an arc of H. Since $a = (u, v)$ is a real arc of H, it follows that u is the right external vertex of the right tied wedge W. Thus $a^{-1} = (v, u)$ is the right external rib of W.

For the second case, suppose a is not an arc of H. Hence (x, v) is a virtual arc of H such that u is the relay vertex of (x, v), (by Definition 2.3.4). This implies there is a wedge W' with relay vertex u such that (x, v) is a virtual arc of W'. Hence a is a rib of W'. However, this contradicts Lemma 2.3.2(iii), since a^{-1} is a rib of W, and $W \neq W'$.
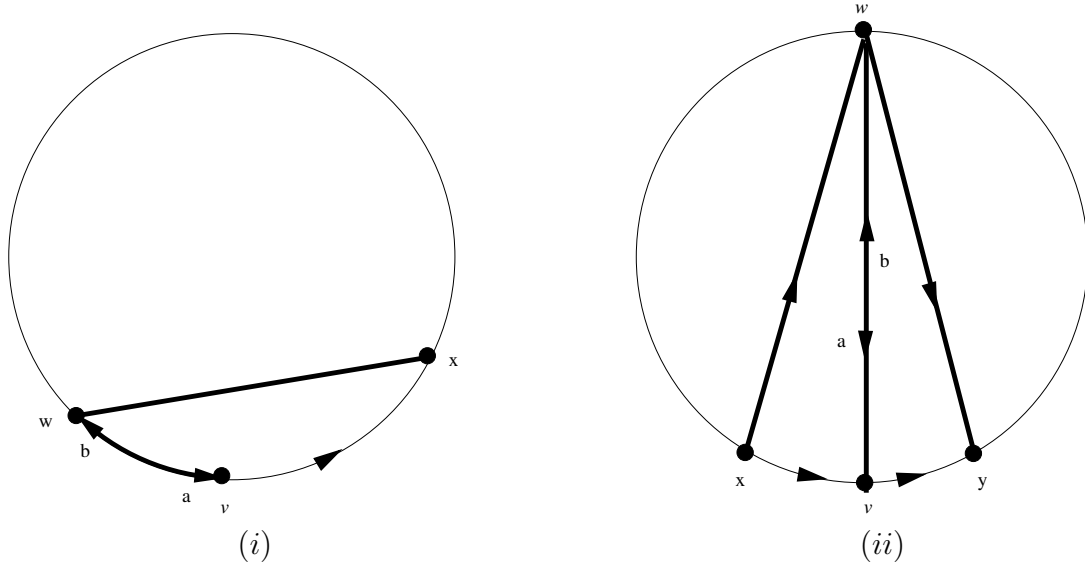
$\square$

In the following, we consider the case when $a = b^{-1}$. The next lemma shows that

in this situation, a and b must be isolated arcs of v.

Lemma 2.3.6. *If a is the incoming arc of v, and b is the outgoing arc of v, and $a = b^{-1}$, then a and b are isolated arcs of v.*

Figure 2.7: *Two of the three possible configurations showing backtrack arcs of v, where a is the incoming arc of v and b is the outgoing arc of v, and $a = b^{-1}$. (i) The vertex v an external vertex of a tied wedge with relay vertex $w \neq v$. (ii) The vertex v is an internal vertex of a wedge with relay vertex $w \neq v$.*

Proof. By Definition 2.3.4, there are two cases for each of a and b, and hence four cases total. First suppose both a and b are arcs of H. Since $a = b^{-1}$, and since H is a directed cycle, H contains only these two arcs, and hence H has two vertices. This is a contradiction since $\vec{G}$ has at least three vertices.

Next suppose a is an arc of H, and b is not. So a is not a rib of any wedge. By Definition 2.3.4, $b = (v, w)$ such that w is the relay vertex of the virtual arc (v, x) of H, see Figure 2.7(i). Now w is the relay vertex for a left tied wedge W, and b is the left external rib of W. By Lemma 2.3.2(ii), b is not a rib of any wedge with relay vertex v. Therefore a and b are isolated arcs of v. If b is an arc of H and a is not, the argument is similar.

Finally, suppose neither a nor b are arcs of H. By Definition 2.3.4, $a = (w, v)$ such that w is the relay vertex of the virtual arc (x, v), and similarly $b = (v, w)$ such that

w is the relay vertex of the virtual arc (v, y). Hence a and b are internal ribs of the wedge W with relay vertex w, see Figure 2.7(ii). By Lemma 2.3.2(ii), this implies neither a nor b are ribs of any wedge other than W. Therefore, they are isolated arcs of v, (since $w \neq v$).

$\square$

We now have a good understanding of how wedges interact with these incoming and outgoing arcs. The next section describes how to assign port numbers to the arcs of $\vec{G}$.

2.3.2 Assigning Port Numbers

In the previous section, we have defined wedges, and identified the incoming and outgoing arcs for every vertex. In this section, we will assign a label to each arc of $\vec{G}$, called a *port number*. Once assigned, these port numbers will be used by the robot in Section 2.1 for the periodic traversal of $\vec{G}$.

Let e be an arc of $\vec{G}$. One port number will be assigned to e, and one will be assigned to e^{-1}. When every arc has a port, we call this a *port numbering* of $\vec{G}$. If each arc has exactly one port number that is chosen uniquely from the set $\{1, 2, 3, \ldots, d(v)\}$, where v is the tail vertex of the arc, we call this a *valid* port numbering of $\vec{G}$. When comparing two port numbers, we use the natural ordering of integers. By associating these port numbers with the tail of the arc, this will induce a local orientation of the underlying graph G, (where each edge has two ports, one at each vertex).

We will now describe how to assign an interval of port numbers to some ribs of a wedge. Given any wedge W with relay vertex v, and virtual path $P = u_1, u_2, \ldots, u_{s(W)}$, we say that an interval $[p, q]$ *is assigned to the wedge* W if $1 \leq p < q \leq d(v)$, $s(W) = q - p + 1$, and port $p + i - 1$ is assigned to (v, u_i) for all $1 \leq i \leq s(W)$. We say that the interval $[p, q]$ *is assigned to* a set of wedges $\{W_1, W_2, \ldots, W_k\} = \{W_i\}_{i=1}^{k}$, if all wedges have a common relay vertex, $\sum_{i=1}^{k} s(W_i) = q - p + 1$, $[p, p + s(W_1) - 1]$ is assigned to W_1, and $[p + \sum_{i=1}^{j-1} s(W_i), p + \sum_{i=1}^{j} s(W_i) - 1]$ is assigned to W_j for $1 < j \leq k$.

The assignment of the port numbers will be accomplished by applying the Port-Numbering procedure detailed below to each vertex v of $\vec{G}$. After identifying the incoming and outgoing arcs of v, there are five ways that these arcs can interact with wedges. This leads to the five exclusive conditions for the what we will refer to as the five subroutines of the Port-Numbering procedure, starting at lines 7, 13, 20, 27, and 34, respectively. Each subroutine will first assign port numbers to the incoming and outgoing arcs of v if they are isolated, and then assigns an interval of ports to every wedge with relay vertex v. The only wedges that are considered seperately by the Port-Numbering procedure are tied wedges. The ports and intervals are chosen so that the incoming and outgoing arcs always have consecutive port numbers. Finally,

the Port-Numbering procedure will run the arbitrary-ports subprocedure, which arbitrarily assigns a port to each arc with tail vertex v that does not already have a port. This procedure also ensures that the port $d(v)$ is not assigned to any arc with tail vertex v by the arbitrary-ports subroutine.

---

procedure Port-Numbering(v)

1: Let a be the incoming arc of v

2: Let b be the outgoing arc of v

3: Set k such that v is a relay vertex of k distinct wedges

4: **if** $k > 0$ **then**

5:    Let $\{W_i\}_{i=1}^k$ be the list of ordered wedges at v

6: **end if**

7: **if** $a = b^{-1}$ **then**

8:    Assign $d(v)$ to b

9:    **if** $k > 0$ **then**

10:      Assign $[d(v) - \sum_i s_i, d(v) - 1]$ to $\{W_i\}_{i=1}^k$

11:    **end if**

12: **end if**

13: **if** $a \neq b^{-1}$ and both a and b are isolated arcs of v **then**

14:    Assign $d(v) - 1$ to a^{-1}

15:    Assign $d(v)$ to b

16:    **if** $k > 0$ **then**

17:      Assign $[d(v) - \sum_i s_i - 1, d(v) - 2]$ to $\{W_i\}_{i=1}^k$

18:    **end if**

19: **end if**

20: **if** $a \neq b^{-1}$ and both a and b are not isolated arcs of v **then**

21:    Assign $[d(v) - s_1 + 1, d(v)]$ to W_1

22:    Assign $[d(v) - s_1 - s_k + 1, d(v) - s_1]$ to W_k

23:    **if** $k \geq 3$ **then**

24:      Assign $[d(v) - \sum_i s_i + 1, d(v) - s_1 - s_k]$ to $\{W_i\}_{1 < i < k}$

25:    **end if**

26: **end if**
27: **if** $a \neq b^{-1}$, a is an isolated arc of v but b is not **then**
28:      Assign $d(v) - s_1$ to a^{-1}
29:      Assign $[d(v) - s_1 + 1, d(v)]$ to W_1
30:      **if** $k \geq 2$ **then**
31:          Assign $[d(v) - \sum_i s_i, d(v) - s_1 - 1]$ to $\{W_i\}_{i>1}$
32:      **end if**
33: **end if**
34: **if** $a \neq b^{-1}$, b is an isolated arc of v but a is not **then**
35:      Assign $[d(v) - s_k, d(v) - 1]$ to W_k
36:      Assign $d(v)$ to b
37:      **if** $k \geq 2$ **then**
38:          Assign $[d(v) - \sum_i s_i, d(v) - s_k - 1]$ to $\{W_i\}_{i<k}$
39:      **end if**
40: **end if**
41: arbitrary-ports(v)

In the following lemma, we show that the Port-Numbering procedure will generate a valid port numbering of $\vec{G}$.

Lemma 2.3.7. *The procedure Port-Numbering will produce a valid port numbering of $\vec{G}$.*

Proof. It is sufficient to show the following:
(i) If a vertex v sees port p, then $1 \leq p \leq d(v)$.
(ii) If two distinct arcs share the same tail vertex, they are assigned different ports.
(iii) Every arc of $\vec{G}$ has been assigned exactly one port.

To show (i), suppose that the arc $e = (v, u)$ has port p. Let a be the incoming arc of v, and b be the outgoing arc of v.

If p was not assigned in an interval, then by inspection of the procedure, p is either $d(v) - 1$, $d(v)$, or $d(v) - s_1$. Consider $p = d(v) - s_1$; the other cases are trivial. So p was assigned on line 28, which is part of the subroutine on line 27. So s_1 is the

size of the wedge W_1. Since b is an isolated arc of v, $2 \leq s_1 \leq d(v) - 1$. Therefore $1 \leq p \leq d(v)$.

Now suppose p belongs to some interval $I = [q, r]$, where I was assigned to some wedge (or set of wedges) with relay vertex v. Let $W_1, W_2, \ldots, W_k$ be the list of ordered wedges at v. Since $q \leq p \leq r$, it is sufficient to show that $q \geq 1$, and $r \leq d(v)$.

First, we show that $q \geq 1$. By inspection of the procedure, q is of the form $d(v) - \sum_i s_i$, $d(v) - \sum_i s_i - 1$, or $d(v) - \sum_i s_i + 1$, where the sum is over some subset of $\{W_i\}_{i=1}^k$. Consider the case where $q = d(v) - \sum_i s_i - 1$ on line 17; the other cases are similar. Since this interval was assigned in the subroutine starting on line 13, both a and b are isolated arcs of v. So $\sum_i s_i \leq d(v) - 2$, which implies $d(v) - 2 - \sum_i s_i \geq 0$, and hence $q \geq 1$.

Next we show $r \leq d(v)$. By inspection of the procedure, r is either $d(v)$, $d(v) - 1$, $d(v) - 2$, $d(v) - s_1$, $d(v) - s_1 - s_k$, or $d(v) - s_i - 1$. However, since the size of a wedge is always positive, this is trivial.

In all cases, $1 \leq p \leq d(v)$, so this proves (i).

To show (ii), let $e = (v, u)$ and $e' = (v, u')$ be two arcs such that $u \neq u'$. Suppose that e was assigned port p, and e' was assigned port p', by the Port-Numbering procedure. If either of p or p' were assigned by the arbitrary-ports subprocedure, we are done since by definition these ports were chosen uniquely. Otherwise, both p and p' were assigned in the same subroutine of the Port-Numbering procedure, (since at most one subroutine is executed for each vertex). Notice that if p and p' are both ports in the same interval, then we are done since $u \neq u'$, (see the definition of how intervals of ports are assigned to wedges). From now on, suppose p and P' are not both in the same assigned interval.

The five subroutines are disjoint, hence we consider them individually. It is sufficient to show, for each subroutine, that every pair of intervals assigned is disjoint, and no port assigned outside of an interval is contained in some assigned interval. This will imply that $p \neq p'$.

Consider first the subroutine on line 7. There is at most one interval assigned, and by line 10, the largest port in that interval is $d(v) - 1$. The port $d(v)$ assigned on line 8 is not contained in this interval. Hence $p \neq p'$ since at most one of p or p' is contained in the interval, by supposition.

Next consider the subroutine on line 13, so the ports $d(v) - 1$ and $d(v)$ were assigned on lines 14 and 15, respectively. There is at most one interval assigned in this subroutine, and by line 17 the largest port contained in this interval is $d(v) - 2$.

Now consider the subroutine on line 20. There were at most three intervals assigned, namely $[d(v) - s_1 + 1, d(v)]$, $[d(v) - s_1 - s_k + 1, d(v) - s_1]$, and $[d(v) - \sum_i s_i + 1, d(v) - s_1 - s_k]$, by lines 21, 22, and 24, respectively. Clearly these are disjoint since $d(v) - s_1 - s_k < d(v) - s_1 - s_k + 1 < d(v) - s_1 < d(v) - s_1 + 1$.

The subroutine on line 27 assigns port $d(v) - s_1$ on line 28, and it assigns at most two intervals. By lines 29 and 31, we can clearly see that these intervals are disjoint and do not contain the port $d(v) - s_1$, since $d(v) - s_1 - 1 < d(v) - s_1 < d(v) - s_1 + 1$.

Finally, consider the subroutine on line 34. So port $d(v)$ was assigned on line 36. There are at most two intervals assigned on lines 35 and 38. These intervals clearly do not contain $d(v)$, and they are disjoint since $d(v) - s_k + 1 < d(v) - s_k$.

Hence all cases, $p \neq p'$, so this proves (ii).

To show (iii), let $e = (v, u)$ be an arc of $\vec{G}$. The Port-Numbering procedure will be executed at v exactly once. Hence the arc e will be assigned a port by at most one of the five subroutines, (starting on lines 7, 13, 20, 27, and 34).

If e is not assigned a port in any of these five subroutines, then by definition that the arbitrary-ports subprocedure will always assign a port to e. So e always has at least one port number.

It remains to show that each subroutine assigns at most one port to e, (since they are exclusive). If e was assigned a port outside of an interval, then e is an isolated arc of v. We know e belongs to at most one wedge by Lemma 2.3.2(ii). It is easy to verify in every case that the length of each interval is equal to the size of the corresponding wedge (or wedges) that the interval was assigned to. For example, the length of the interval on line 10 is $d(v) - 1 - (d(v) - \sum_i s_i) + 1 = \sum_i s_i$.

Finally, we show that each wedge is assigned at most one interval. The only case where this may not be clear is in the subroutine on line 20. Let a be the incoming arc of v, b be the outgoing arc of v, and let $W_1, W_2, \ldots, W_k$ be the list of ordered wedges at v. By Lemma 2.3.5, a^{-1} is a rib of the tied wedge W_1 and b^{-1} is a rib of the tied wedge W_k. We also know by Lemma 2.3.3 that W_1 is not right tied and W_k is not left tied, so $k \neq 1$. Therefore each wedge is assigned at most one interval. and e has

exactly one port number.

$\square$

Now we know that the Port-Numbering procedure is well defined, and it produces a valid port numbering of the graph $\vec{G}$. In the next section, we define a robot that will perform the traversal of $\vec{G}$ by using these ports.

2.3.3 The Invariant

Now we define an invariant which will guarantee, under certain conditions, that the robot will periodically traverse $\vec{G}$. Let H^* be the directed graph constructed by starting from H, then replacing every virtual arc of H with its corresponding relay path. Naturally, H^* induces a (directed) Hamiltonian walk of $\vec{G}$ by following H, where each virtual arc is traversed via the corresponding relay path. Hence H^* is a subgraph of $\vec{G}$. We will abuse notation and identify the directed graph H^* with this Hamiltonian walk of $\vec{G}$, and denote both by H^*. When we refer to a port number of an arc of H^*, we of course mean the port number of the corresponding arc of the graph $\vec{G}$. Observe that if W is a wedge, then every rib of W is in H^*. This implies that the incoming and outgoing arcs of a vertex are always in H^*.

The fact that H^* is a Hamiltonian walk gives rise to the natural successor function defined on arcs of H^*; the *successor* of an arc e of H^* is the arc following e on H^*, denoted by e^+. In particular, if e is the incoming arc of v, then e^+ is the outgoing arc of v, see for example Figure 2.6. With this in mind, we define an invariant which will guarantee, under certain conditions, that the robot will periodically traverse $\vec{G}$.

The Invariant: *If the robot enters a vertex on an arc e of H^*, then the robot will leave that vertex on the arc e^+ of H^*.*

The invariant guarantees a periodic traversal once the robot uses an arc of H^* to enter a vertex. The fact that the invariant always holds follows from Lemma 2.3.9. First, we first prove that, given the correct initial configuration with any initial vertex, the robot always leaves the vertex on an arc of H^*, and hence enters the new current vertex on an arc of H^*.

Lemma 2.3.8. *Let v be any vertex of G. Suppose the robot starts on the configuration $(d(v), v, \alpha)$. Then the robot will leave v on an arc e of H^*. Hence if $e = (v, u)$, the robot will enter u on an arc of H^*.*

Proof. Let a be the incoming arc of v, b be the outgoing arc of v, and suppose that v is the relay vertex of $k \geq 0$ distinct wedges. Since the configuration of the robot has the entering port $d(v)$, the second case of the transition function will be applied.
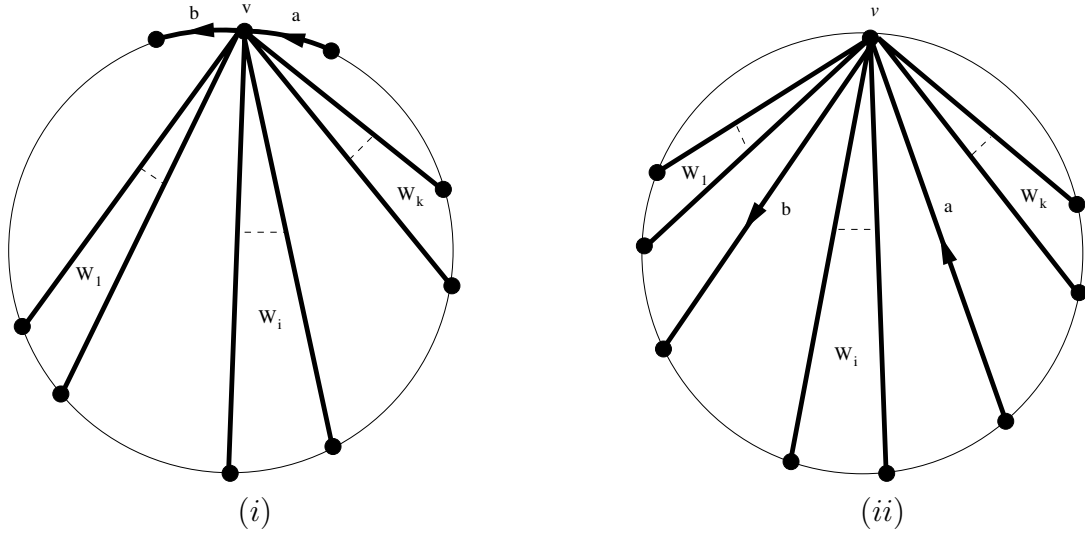
Figure 2.8: *Both a and b are isolated arcs of v. Represented are two of the four possible configurations, depending on whether or not a or b are arcs of H. In (i) we depict the case where both a and b are arcs of H; in (ii) we depict the case where neither a nor b are arcs of H.*

Hence it is sufficient to show that the Port-Numbering procedure has assigned port $d(v)$ to an arc with tail vertex v, that is an arc of H^*. Recall that H^* contains both a and b, as well as any arc that is a rib of any wedge. There are five cases based on the subroutines of the Port-Numbering procedure.

Case 1: Suppose first that $a = b^{-1}$. So the subroutine on line 7 was executed v. Port $d(v)$ was assigned to b by line 8, which is an arc of H^*.

From now on, suppose $a \neq b^{-1}$.

Case 2: Suppose that both a and b are isolated arcs of v; for example see Figure 2.8. So the subroutine on line 13 was executed at v. So b has port $d(v)$ by line 15.

Case 3: Now suppose neither a nor b are isolated arcs of v; see Figure 2.10. So the subroutine at line 20 was executed at v. By line 21, the right external rib of W_1, (which is an arc of H^*), has port $d(v)$.

Case 4: Now suppose b is not an isolated arc of v, and a is an isolated arc of v, for the two possible cases see Figure 2.9. So the subroutine at line 27 was executed at v. Hence the right external rib of W_1, (which is an arc of H^*), was assigned port $d(v)$ by line 29.

Case 5: Finally suppose a is not an isolated arc of v, and b is an isolated arc of v, so the subroutine at line 34 was executed at v. For the two possible cases see Figure 2.11. By line 36, b has port $d(v)$.

$\square$
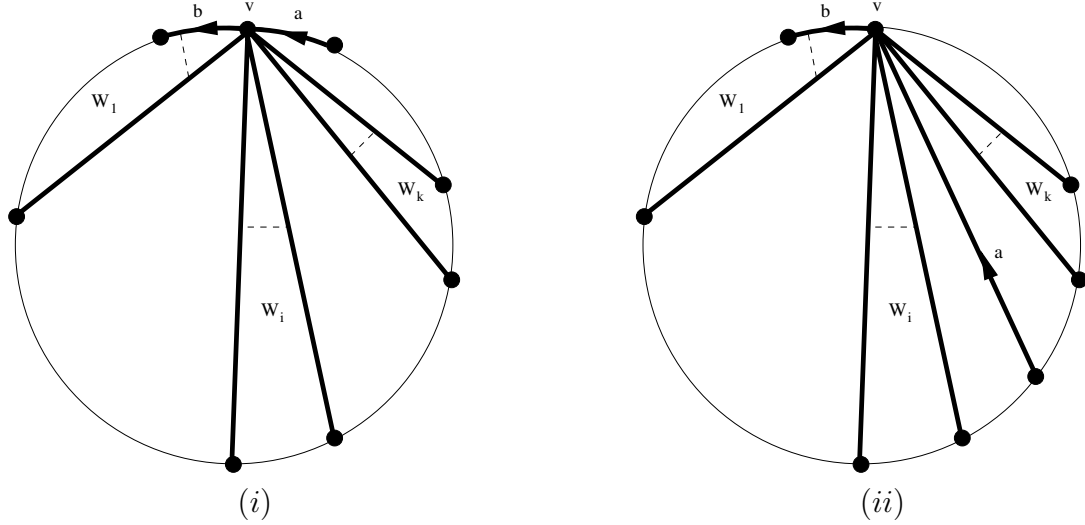
Figure 2.9: *The arc a is an isolated arc of v, and b is not. In (i), a is an arc of H, and in (ii) it is not.*

Now we know that at the initial step, the robot always enters the new current vertex along an arc of H^*. The following lemma implies that the invariant always holds, provided that the robot entered a vertex v on an arc of H^*.

Lemma 2.3.9. *Suppose the current configuration of the robot is (p, v, α), such that v was entered on an arc e of H^*. Then the robot will leave on e^+, (which is an arc of H^*).*

Proof. Let $e = (u, v)$, hence p is the port number of $e^{-1} = (v, u)$. Suppose $T(p, v, \alpha) = (q, v', \alpha)$, so the robot leaves v on the arc (v, v'), and the port of (v', v) is q. It is sufficient to show that $e^+ = (v, v')$.

Let a be the incoming arc of v, b be the outgoing arc of v, and v be the relay vertex of $k \geq 0$ wedges. First notice that, since the tail of b is v, $e \neq b$, and similarly e is not the right external rib of any wedge with relay vertex v. Since e is an arc of

H^*, either $e = a$, or e is an internal rib, or the left external rib of some wedge with relay vertex v.

In what follows, we will examine all possible values for the port p. This will be done with respect to each of the five subroutines of the Port-Numbering procedure, (since the subroutines are exclusive, and we know p was assigned in one of these subroutines since e is an arc of H^*).
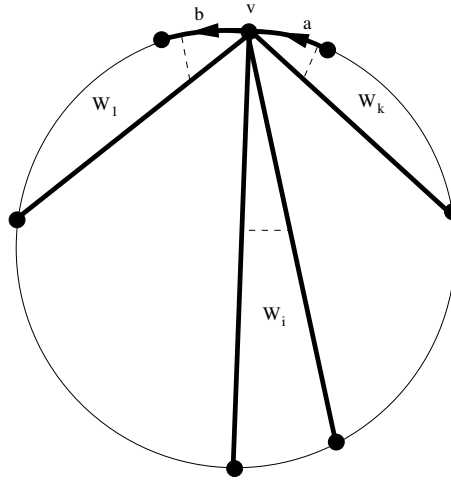

Figure 2.10: *Neither a nor b are isolated arcs of v.*

Case 1: Suppose first that $a = b^{-1}$; see Figure 2.7. So the subroutine on line 7 was executed at v. By line 8, the port of $a^{-1} = b$ is $d(v)$. Hence if $p = d(v)$, then $e = a$, which implies $e^+ = b$. Since the port of e^+ is $d(v)$, $e^+ = (v, v')$.

Now let $p < d(v)$. Since the port of a^{-1} is $d(v)$, $e \neq a$. Since e is an arc of H^*, and $e \neq a$, it follows that e is a rib of some wedge with relay vertex v, and hence $k > 0$. Let $W_1, W_2, \ldots, W_k$ be the list of ordered wedges at v. So e is a rib W_j for some $1 \leq j \leq k$. Furthermore, since e is not the right external rib of W_j, there is a virtual arc (u, x) of W_j such that $e^+ = (v, x)$ is a rib of W_j. Clearly, e^+ was assigned port $p + 1$ at v by line 10. Since $p < d(v)$, $x = v'$ and hence $e^+ = (v, v')$.

From now on, suppose $a \neq b^{-1}$.

Case 2: Suppose both a and b are isolated arcs of v, see Figure 2.8. So the subroutine on line 13 was executed at v. If $p = d(v) - 1$, then $e = a$ by line 14, which implies $e^+ = b$. Since b has port $d(v)$ by line 15, $e^+ = (v, v')$.
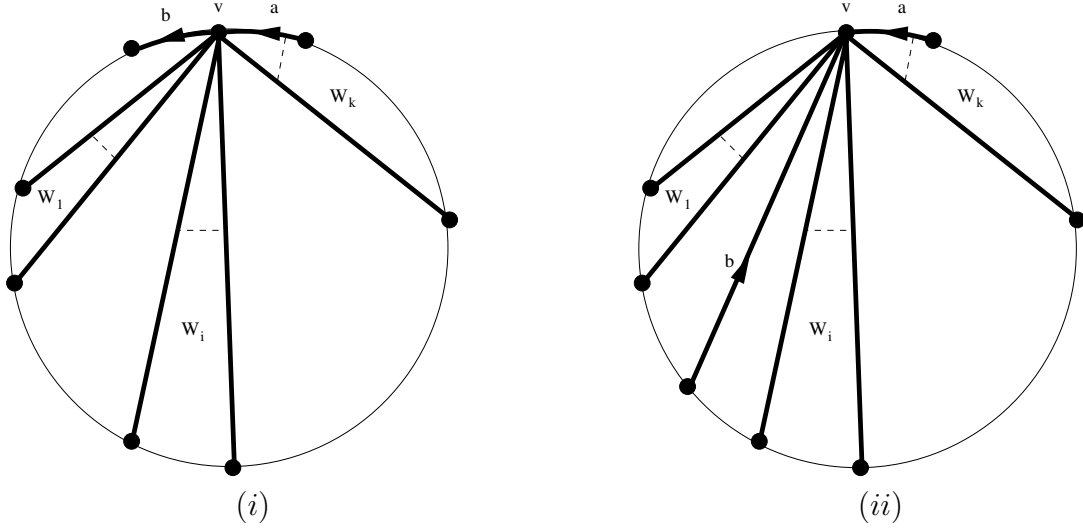
Figure 2.11: *The arc b is an isolated arc of v, and a is not. In figure (i), b is an arc of H, and in figure (ii) it is not.*

Now let $p \neq d(v) - 1$. This implies $e \neq a$, and hence e is a rib of some wedge with relay vertex v, which means $k > 0$. Let $W_1, W_2, \ldots, W_k$ be the list of ordered wedges at v. Let e is a rib of W_j for some $1 \leq j \leq k$. Since the head of e is v, e is either an internal rib or the left external rib of W_j. So $e^+ = (v, x)$ where (u, x) is a virtual arc of W_j. We know by by line 17 that e^+ has port $p + 1$, and $p < d(v)$. Therefore $x = v'$, so $e^+ = (v, v')$.

Case 3:   Suppose that neither a nor b are isolated arcs of v, see Figure 2.10. This implies that the subroutine on line 20 was executed at v. This also implies that $k \geq 1$, so let $W_1, W_2, \ldots, W_k$ be the list of ordered wedges at v. Hence, by Lemma 2.3.5, a^{-1} is the right external rib of W_j, and b^{-1} is the left external rib of $W_{j'}$, for $1 \leq j, j' \leq k$. Furthermore, since W_j is right tied and $W_{j'}$ is right tied, $j = k$ and $j' = 1$. Hence b has port $d(v) - s_1 + 1$, and a^{-1} has port $d(v) - s_1$, by lines 21 and 22, respectively.

Suppose $p = d(v) - s_1$, so $e = a$ and $e^+ = b$. Since e^+ has port $p+1 = d(v) - s_1 + 1$, and since $p < d(v)$, $e^+ = (v, v')$.

Next suppose $p \neq d(v) - s_1$, which implies $e \neq a$. Thus e is either an internal rib or the left external rib of some wedge W_j for $1 \leq j \leq k$. Thus v is the relay vertex of a virtual arc (u, x) of W_j, such that $e^+ = (v, x)$. So port $p + 1$ was assigned to e^+ (by either line 21, 22, or 24). Therefore $x = v'$, hence $e^+ = (v, v')$.

Case 4: Next, we suppose that b is not an isolated arc of v, and a is an isolated arc of v, see Figure 2.9. So the subroutine at line 27 was executed at v. Since b is not an isolated arc of v, we have $k > 0$, so let $W_1, W_2, \ldots, W_k$ be the list of ordered wedges at v. By Lemma 2.3.5, b^{-1} is the left external rib of W_j for some $1 \leq j \leq k$. Since W_j is a left tied wedge, $j = 1$.

If $p = d(v) - s_1$, then $e = a$, which implies $e^+ = b$. By line 29, e^+ has port $d(v) - s_j + 1 = p + 1$, so $e^+ = (v, v')$.

Otherwise $p \neq d(v) - s_1$, so $e \neq a$. Hence e is a rib of some wedge W_j with relay vertex v, where $1 \leq j \leq k$. Since e is not the right external rib of W_j, we have $p < d(v)$. As above, e^+ is a rib of W_j such that e^+ has port $p + 1$, either by line 29 or 31. Therefore $e^+ = (v, v')$.

Case 5: Now suppose that b is an isolated arc of v and a is not, see Figure 2.11. So the subroutine on line 34 was executed at v. Since a is not an isolated arc of v, $k > 0$, so let $W_1, W_2, \ldots, W_k$ be the list of ordered wedges at v. By Lemma 2.3.5, a^{-1} is the right external rib of W_k. By line 35, a^{-1} has port $d(v) - 1$.

If $p = d(v) - 1$, then $e = a$, which implies $e^+ = b$. Since e^+ has port $p + 1 = d(v)$ by line 36, $e^+ = (v, v')$.

Now suppose $p \neq d(v) - 1$. So e is either an internal rib, or the left external rib of W_j, for some $1 \leq j \leq k$. As in previous cases, e^+ is the rib of W_j that has port $p + 1$ by either line 35 or 38, (and it is clear that $p < d(v)$). Thus $e^+ = (v, v')$.

In all five cases, the robot left v on the arc e^+ according to the transition function. The proof of the invariant follows.

$\square$

2.4 Upper Bound of the Period

We obtain an upper bound on the period of this traversal.

Theorem 2.4.1. *Let G be a P_3-free graph, with $\delta(G) \geq 2$. Then there exists a local orientation of G and a corresponding robot (defined in Section 2.1) that will perform a periodic traversal of G with period $\pi(n) \leq 2n - 2$, (provided the robot starts in configuration $(d(v), v, \alpha)$, where v is any vertex of G).*

Proof. If G is not 2-connected, then by Lemma 2.0.1, $\Delta(G) = n - 1$. Hence we will fix the local orientation of G using the Star-Numbering procedure. By Lemma 2.2.1, every vertex of G is visited periodically, with period at most $2n - 3$.

Now suppose G is 2-connected. The result of Fleishner from [15] guarantees that H exists. Suppose H^*, the robot, and the ports of G are as described above. By Lemma 2.3.8, the robot will leave v on an arc of H^*. After this step, by Lemma 2.3.9, the robot will start periodically traversing arcs of H^*, and hence every vertex of G.

Next we estimate the period of the traversal, $\pi(n)$. Recall that the period $\pi(n)$ of the traversal is the maximum number of arcs traversed between two consecutive visits of any vertex.

For any arc $e = (u, v)$ of H, the robot will traverse at most two arcs to get from u to v; one if e is real and two if e is virtual.

The maximum number of arc traversals between two visits of v will clearly occur when the robot traverses the entire walk H^* after leaving v. Suppose that H has i virtual arcs, and j real arcs, where necessarily $i + j = n$. For any arc $e = (u, v)$ of H, the robot will traverse at most two arcs to get from u to v; one if e is real and two if e is virtual. Thus the maximum number of arcs of H^* traversed by the robot is $2i + j$. It was shown in [18], (which is a short constructive proof of Fleishner's theorem), that for every 2-connected graph G, G^2 has a Hamiltonian cycle that contains at least two edges of G. Since H was chosen to have the minimal number of virtual arcs, this result implies that $j \geq 2$. So we maximize $2i - j$ by setting $i = n - 2$, to obtain $\pi(n) \leq 2(n - 2) + 2 = 2n - 2$.

Therefore, in both cases, $\pi(n) \leq 2n - 2$.

$\square$

2.5   Conjectures and Future Work

This chapter provided a method to periodically traverse any P_3-free graph G with $\delta(G) \geq 2$. We conjecture that the techniques given in this chapter can be modified slightly to apply to all graphs.

By increasing the number of states of the robot, and assigning port numbers more carefully in certain cases, we conjecture that we can drop the $\delta(G) \geq 2$ condition, and traverse any P_3-free graph.

Chapter 3

Graph Traversal Using Walks

In this section, we let G be a finite graph with n vertices and m edges that is not anonymous, i.e., every vertex has a unique identifier drawn from set of labels. For example, each vertex could be assigned a unique number from the set $\{1, 2, \ldots, n\}$. Hence an algorithm on such a graph that uses these labels must use at least $O(\log n)$ bits of memory. Recall that a solution to the graph traversal problem is an algorithm that will start at any vertex of G, and will eventually halt after every vertex of G has been visited. Notice how this differs from a periodic traversal, where each vertex must be visited infinitely many times in a periodic manner.

In this section, we provide a traversal algorithm that can start with any vertex of the graph. Our solution is an algorithm that visits edges instead of vertices, and will output a list containing each edge exactly once. Edges are reported to this list in some order as they are visited, (not neccessarily the first time), which implies every vertex has been visited. The traversal is based on two concepts of the graph which we will define, called an oriented walk cover, and a compatible total order. These concepts are inspired by faces of a planar graph, and the geometric information of edges and vertices in an embedding of the graph. Then, we will define an auxilliary graph called the virtual tree, which is used to prove the correctness of our algorithm. We show that our algorithm will traverse any graph G with these two concepts in $O(m \log m)$ time and $O(\log n)$ memory, where G has m edges and n vertices.

3.1 Covering Walks and a Compatible Total Order

Let G be an undirected, finite, simple graph, and let $\vec{G}$ be its symmetric orientation. Recall that $\vec{G}$ replaces each edge uv in G by the pair of arcs (u, v) and (v, u), and that we refer to the reverse of $e = (u, v)$ by $e^{-1} = (v, u)$. The following definition is inspired by the idea of faces in a planar graph. It is also similar to a cycle double cover, which we discuss in Section 3.4.

Definition 3.1.1. *An oriented walk cover W of $\vec{G}$ is a collection of closed walks from $\vec{G}$ such that each arc e of $\vec{G}$ is in exactly one walk $w \in W$. Furthermore, for any arc e of $\vec{G}$, if $e \in w$, then $e^{-1} \notin w$. We call w a covering walk of $\vec{G}$.*

It is clear that any oriented walk cover of $\vec{G}$ is finite, since G is finite. We say the size of a covering walk w, denoted $|w|$, is the number of arcs of w. Next we define a special total order on the arcs of $\vec{G}$ such that any arc and its reverse are consecutive in the total order.

Definition 3.1.2. *Let G be an undirected graph, and suppose $\vec{G}$ has a total order on its arcs. We say this total order is compatible if for any arc e of $\vec{G}$, there does not exist an arc e' such that either $e < e' < e^{-1}$ or $e^{-1} < e' < e$.*

From now on, suppose $\vec{G}$ has a compatible total order, and an oriented walk cover W. Every total order over a finite set has a unique minimal element, hence we denote the minimal arc of $\vec{G}$ by e_0. Furthermore, for every covering walk w of $\vec{G}$, there is a minimal arc with respect to this total order restricted to the arcs of the walk w. We call this arc the *entering arc* of w.

We will now proceed to define several functions that use $O(\log n)$ memory. For any arc e, we say **cwalk**$(e) = w$ if w is the unique covering walk containing e. If $e = (u, v)$, let **rev**$(e) = e^{-1}$, **head**$(e) = v$, and **tail**$(e) = u$. Next define the predecessor function **pred**$(e) = f$ if and only if f is in **cwalk**(e) and **head**$(f) = $ **tail**(e). Similarly, the successor function **succ**$(e) = f$ if and only if f is in **cwalk**(e) and **tail**$(f) = $ **head**(e).

Given any covering walk w we say **entry**$(w) = e$ if e is the entering arc of w. We compute this using the algorithm from [4], which was also used in [6]. For convenience we give this algorithm below. Recall that our traversal algorithm uses a robot that visits one arc of $\vec{G}$ at a time. Initially, the current arc of the algorithm is e, and

the algorithm makes two copies of the robot. These two robots move clockwise and counter-clockwise, respectively. The algorithm halts whenever a smaller arc is found, or when the two robots run into each other, at which point the traversal is resumed with the original robot at e, (and the cloned robots can be removed).

Algorithm entry$(w) = e$

1: $e^{cw} \leftarrow e$
2: $e^{ccw} \leftarrow e$
3: **repeat**
4:      $e^{cw} \leftarrow \mathbf{pred}(e^{cw})$
5:      **if** $e^{cw} = e^{ccw}$ **then**
6:          **return** true
7:      **else**
8:          **if** $e^{cw} < e$ **then**
9:              **return** false
10:         **end if**
11:     **end if**
12:     $e^{ccw} \leftarrow \mathbf{succ}(e^{ccw})$
13:     **if** $e^{cw} = e^{ccw}$ **then**
14:         **return** true
15:     **else**
16:         **if** $e^{ccw} < e$ **then**
17:             **return** false
18:         **end if**
19:     **end if**
20: **until** true

This algorithm will always halt, since w is a closed walk. The correctness of the **entry** is easy to check, since it only returns true after it has compared the given arc

e to every other arc of the walk. If **entry**(e) returns false, then there is some arc e' such that $e' < e$. Since only three robots are needed, we can say this algorithm uses a constant number of mark bits, (or pebbles), hence at most $O(\log n)$ memory is required. The purpose of using this method is that it has sublinear running time, as opposed to simply traversing the walk in only one direction which takes $O(n)$. The running time of this algorithm follows from the following Lemma from [4], which we repeat here along with the proof. This lemma was originally given in terms of faces of a planar graph, instead of closed walks, however the result still holds.

Lemma 3.1.3. [4] *Let $e_1, e_2, \ldots, e_{|w|}$ be the arcs of the covering walk w in order. We say that e_i is k-minimum if $e_i \leq e_j$ for all $i - k \leq j \leq i + k$, where subscripts are taken $\mathrm{mod}|w|$. We define $mk(e_i)$ as the maximum k for which e_i is k-minimum. Then*
$$\sum_{i=1}^{|w|} mk(e_i) \leq |w| \ln |w|.$$

Proof. If e_i is k-minimum, then none of $e_{i-k}, \ldots, e_{i-1}, e_{i+1}, \ldots, e_{i+k}$ are k-minimum. Therefore, at most $\lfloor |w|/(k+1) \rfloor$ arcs of w are k-minimum. Thus,

$$\sum_{i=0}^{|w|-1} mk(e_i) = \sum_{k=1}^{|w|-1} |\{e_i \ : \ e_i \text{ is } k\text{-minimum}\}| \leq |w| \sum_{k=1}^{|w|-1} \lfloor 1/(k+1) \rfloor \leq |w| \ln |w|,$$

where the final step follows since $\sum_{i=1}^{x} 1/x \leq \ln x + 1$.

$\square$

The traversal algorithm will start by traversing some current covering walk. Next, we define a relationship between certain covering walks. We define **parent**$(w) =$ **cwalk**(**rev**(**entry**(w))), provided that $e_0 \neq$ **entry**(w). Hence, if e is the entering arc of w, and $e \neq e_0$, then e^{-1} is an arc of the covering walk **parent**(w). The idea of the next lemma is to show that this parent function is decreasing, i.e., it produces a walk with a smaller entering arc.

Lemma 3.1.4. *If $w \in W$ such that $e_0 \neq$ **entry**(w), and we suppose* ***entry**(**parent**(w)) $\neq$ **rev**(**entry**(w)), then **entry**(**parent**(w)) $<$ **entry**(w).*

Proof. Let $e_w =$ **entry**(w) and $e_p =$ **entry**(**parent**(w)). Notice that since e_p and e_w are in different walks, $e_p \neq e_w$, and $e_p \neq e_w^{-1}$ by supposition.

Suppose to the contrary that $e_w < e_p$. We know $e_p < e_w^{-1}$, since e_p is the entering arc of **parent**(w). But now $e_w < e_p < e_w^{-1}$, a contradiction since we supposed the total order is compatible. $\qquad\square$

Next we define a function that is able to determine whether any given arc is the minimal arc e_0, without comparing the given arc to every other arc of $\vec{G}$. Define function **ismin**(e) to return true if and only if the following three conditions are satisfied:

- $e = \mathbf{entry}(\mathbf{cwalk}(e))$

- $e^{-1} = \mathbf{entry}(\mathbf{cwalk}(e^{-1}))$

- $e < e^{-1}$

Next we prove the correctness of the **ismin**() function.

Lemma 3.1.5. *The function **ismin**(e) returns true if and only if $e = e_0$.*

Proof. Let e be an arc of $\vec{G}$. Suppose **ismin**(e) returns true, i.e., its three conditions are satisfied. Suppose to the contrary that $e_0 \neq e$, which implies $e_0 < e$. Let $w = \mathbf{cwalk}(e)$, and $e_p = \mathbf{entry}(\mathbf{parent}(w))$. By definition, $\mathbf{parent}(w) = \mathbf{cwalk}(\mathbf{rev}(\mathbf{entry}(w))) = \mathbf{cwalk}(e^{-1})$. This implies $e_p = \mathbf{entry}(\mathbf{cwalk}(e^{-1})) = e^{-1}$, by the second condition of **ismin**(e). Since $e = \mathbf{entry}(w)$, it follows from Lemma 3.1.4 that $e_p = e^{-1} < e$. This is a contradiction, since we supposed $e < e^{-1}$.

Now we suppose $e = e_0$ and show **ismin**(e) must return true. It is clear that $e < e^{-1}$ and $e = \mathbf{entry}(\mathbf{cwalk}(e))$. Let $e' = \mathbf{entry}(\mathbf{cwalk}(e^{-1}))$. Suppose to the contrary that $e' \neq e^{-1}$, which implies $e' < e^{-1}$. But now $e < e' < e^{-1}$, which is a contradiction since this is a compatible total order. $\qquad\square$

3.2   Virtual Tree

The main idea of the traversal algorithm is to traverse every covering walk, and switch between covering walks using these entering arcs. Now we define a virtual tree, which is inspired by the idea of the dual of a planar graph. This virtual tree will use covering walks as vertices, and consider two walks to be adjacent whenever one walk is the parent of the other.

Given the oriented walk cover W of $\vec{G}$, we construct an auxiliary graph $G(W)$ called the *virtual tree*. Define $G(W) = (W, E(W))$ where $E(W) = \{(w, w') : \mathbf{parent}(w) = w'\}$. We will consider the walk containing e_0 to be the root of the virtual tree. We next show that $G(W)$ is a tree by showing it is acyclic and connected.

Lemma 3.2.1. *$G(W)$ is a tree with root $\mathbf{cwalk}(e_0)$.*

Proof. Let $w \in W$, and let $w_0 = \mathbf{cwalk}(e_0)$. It is sufficient to show there is a unique path in $G(W)$ from w to w_0. First, notice that there cannot be more than one path from w to w_0, since each walk has exactly one entering arc.

Suppose to the contrary that there is no path from w to w_0. Consider the sequence $w_1 = \mathbf{parent}(w)$, $w_2 = \mathbf{parent}(\mathbf{parent}(w))$, and so on. So $w_i \neq w_0$ for all $i > 0$ by assumption. Since $\mathbf{entry}(w_0) = e_0$, we have $\mathbf{entry}(w_i) > \mathbf{entry}(w_0) = e_0$ for all $i > 0$.

We will use Lemma 3.1.4, which means we must show $\mathbf{entry}(\mathbf{parent}(w_i))$ and $\mathbf{rev}(\mathbf{entry}(w_i))$ are not equal for all $i > 0$. Suppose to the contrary that $\mathbf{entry}(\mathbf{parent}(w_i)) = \mathbf{rev}(\mathbf{entry}(w_i))$ for some i. Let $e = \mathbf{entry}(w_i)$, so by assumption $e^{-1} = \mathbf{entry}(\mathbf{parent}(w_i))$. Therefore the first two conditions of $\mathbf{ismin}()$ are satisfied. Since $e \neq e^{-1}$, we have either $e < e^{-1}$ or $e^{-1} < e$. Therefore either $\mathbf{ismin}(e)$ or $\mathbf{ismin}(e^{-1})$ returns true. This contradicts 3.1.5, since none of w_i, in particular $\mathbf{cwalk}(e)$ or $\mathbf{cwalk}(e^{-1})$, are w_0. We conclude that $\mathbf{entry}(\mathbf{parent}(w_i)) \neq \mathbf{rev}(\mathbf{entry}(w_i))$, so by Lemma 3.1.4, $\mathbf{entry}(w_{i+1}) < \mathbf{entry}(w_i)$ for all $i > 0$. Hence, the infinite sequence $w_1, w_2, w_3, \ldots$ is a strictly decreasing sequence that is bounded below by $\mathbf{entry}(w_0)$. This is a contradiction since $\vec{G}$ and W are finite. Therefore there is exactly one path in $G(W)$ from w to w_0, and hence $G(W)$ is a tree with root $\mathbf{cwalk}(e_0)$. $\square$

The traversal algorithm is given in the next section.

3.3   Traversal Algorithm

The traversal algorithm, given below, will start on any given arc e as the current arc. First, it will find e_0. Then, it will set the current arc to e_0^{-1}, and proceed to move around this current walk. It will move to a new covering walk whenever the reverse of the current arc is the entering arc for that walk. In this way, it simulates a Depth-First search of the virtual tree. It will halt when it again reaches e_0. Even though vertices and arcs of $\vec{G}$ may be visited multiple times, the algorithm reports each edge of G exactly once.

Algorithm Walk-Traversal(e)

1: **repeat** {find the minimum arc e_0}
2:    $e \leftarrow \mathbf{rev}(e)$
3:    **while** $e \neq \mathbf{entry}(\mathbf{cwalk}(e))$ **do**
4:       $e \leftarrow \mathbf{succ}(e)$
5:    **end while**
6: **until ismin**(e) =true
7: **repeat** {start the traversal}
8:    $e \leftarrow \mathbf{succ}(e)$
9:    **if entry**$(\mathbf{cwalk}(e)) < \mathbf{entry}(\mathbf{cwalk}(\mathbf{rev}(e)))$ **then** {check to report}
10:       **report** e
11:    **end if**
12:    **if** $e = \mathbf{entry}(\mathbf{cwalk}(e))$ **then** {finished with current walk}
13:       $e \leftarrow \mathbf{rev}(e)$
14:    **else**
15:       **if rev**$(e) = \mathbf{entry}(\mathbf{cwalk}(\mathbf{rev}(e)))$ **then** {switch to adjacent walk}
16:          $e \leftarrow \mathbf{rev}(e)$
17:       **end if**
18:    **end if**
19: **until ismin**(e) =true

Next we prove the correctness of the Walk-Traversal algorithm, as well as showing its memory and running time.

Theorem 3.3.1. *Let G be a graph with n vertices and m edges. Suppose that G has an oriented walk cover, and $\vec{G}$ has a compatible total order on its arcs. Then the Walk-Traversal algorithm reports every edge of G exactly once using $O(m \log m)$ time and $O(\log n)$ memory.*

Proof. The fact that the algorithm uses $O(\log n)$ memory is clear, since every function used in the Walk-Traversal algorithm uses no extra memory, and they are used in a sequential order (i.e., not recursively). The proof of correctness, and of the running time of the Walk-Traversal algorithm, are similar to the proofs in [4]. Next, we show the Walk-Traversal algorithm simulates a DFS of $G(W)$, and that each edge of G is reported exactly once.

The algorithm starts with an arc e in $\vec{G}$ as input. First, starting on line 3, the algorithm searches for e_0 by starting with **cwalk**(e^{-1}); if e_0 is not found in this walk, it moves to the parent walk and the search is repeated. Since G is finite, and by Lemma 3.1.4, this process will always halt at e_0. Let $w = $ **cwalk**(e_0). Next, the algorithm moves to $e = $ **succ**(e_0) on line 8. If e^{-1} is the entering arc of its covering walk $w' = $ **cwalk**(e^{-1}), the algorithm moves to e^{-1} on line 16. Then the algorithm begins to traverse w' using the **succ**() function. Eventually, (since w' is a closed walk), the algorithm will visit e^{-1} again, (after entering and returning from every covering walk that is under w' in the tree $G(W)$). Since e^{-1} is the entering arc of w', the algorithm moves to e by line 13, and continues traversing the covering walk w using the **succ**() function. The algorithm halts when it returns to e_0 again. Hence it is clear that this simulates a depth first search of $G(W)$. Furthermore, for each arc e, exactly one of e or e^{-1} is reported by line 10, (since they are in different walks, which in turn cannot have the same entering arc). This implies each edge of G is reported exactly once.

Now we compute the running time of the Walk-Traversal algorithm. Ignoring the computation of **entry**() for now, each arc is visited once in the initial phase, (while searching for e_0), and once during the DFS traversal. So far, this implies $O(m)$ arcs have been visited. Next, we will consider **entry**(), so let w be a covering walk. If we

compute $\mathbf{entry}(w) = e$ for each arc e of w, by Lemma 3.1.3 at most $2|w|\log|w|$ arcs are visited. By inspection of the Walk-Traversal algorithm, the $\mathbf{entry}()$ function is executed at most 9 times for each arc e of G, (four from two executions of $\mathbf{ismin}()$, twice when the algorithm is traversing $\mathbf{cwalk}(e^{-1})$, plus three others). So algorithm visits at most $18\sum_{w\in W}|w|\log|w|$ arcs due to the $\mathbf{entry}()$ function. Since $\sum_{w\in W}|w| = O(m)$, and $|w| \leq m$, we obtain $18\sum_{w\in W}|w|\log|w| = O(m\log m)$. Therefore the running time of the Walk-Traversal algorithm is $O(m\log m) + O(m) = O(m\log m)$.

$\square$

3.4 Conjectures and Future Work

The algorithm in the previous section applies to any graph G that has two properties: an oriented walk cover of $\vec{G}$, and a compatible total order on the arcs of $\vec{G}$. In this section we conjecture which graphs have these properties, based on certain assumptions. First we define a cycle double cover of an undirected graph.

Definition 3.4.1. *A cycle double cover C of a graph G is a collection of cycles of G such that each edge of G is in exactly two cycles from C.*

The *cycle double cover conjecture* was posed independently in [24] and [25]; it states that every graph that does not contain a cut edge has a cycle double cover. There is a similar conjecture about directed cycles. The *oriented cycle double cover conjecture*, if it is true, implies that if G is a graph with no cut edge, then there is a collection of cycles in $\vec{G}$ such that each arc is in exactly one cycle [19]. Hence any graph that satisfies this conjecture has an oriented walk cover. For example, the results in [20] imply that G has a cycle double cover if G is either 4-edge connected, or if G cubic and 3-colorable.

Next, we investigate a similar conjecture related to graph embeddings. In the following, we only consider orientable surfaces. First we define a specific type of embedding of a graph.

Definition 3.4.2. *An embedding of a graph G on a surface Σ is called a circular embedding every face of the embedding is a cycle of G.*

There is an open problem, called the *circular embedding conjecture* [19], stating that every 2-connected graph has a circular embedding. It is clear that this conjecture will imply that $\vec{G}$ has an oriented walk cover. Notice that if a graph has a circular embedding, it has a cycle double cover.

Next, we define an ordering on the arcs of $\vec{G}$, based on a specific geometric embedding. We conjecture that this is a compatible total order. If u is a vertex of G, then we will denote the cartesian coordinates of u in the embedding by $(\mathbf{x}(u), \mathbf{y}(u), \mathbf{z}(u))$. Also, we say that, if $e = uv$, **smaller**$(e) = u$ whenever $(\mathbf{x}(u), \mathbf{y}(u), \mathbf{z}(u)) < (\mathbf{x}(v), \mathbf{y}(v), \mathbf{z}(v))$, using lexographic comparison. We will now define an identifier for each arc $e = (u, v)$ of $\vec{G}$, denoted **key**(e).

We will require that these coordinates are bounded rational numbers to define our ordering. The following theorem shows that in fact there is an embedding for any graph such that these coordinates are always integers, called a grid embedding.

Theorem 3.4.3. [8] *Let G be a graph with n vertices. Then there is a three dimensional grid embedding of G such that no pair of edges intersects, and the coordinates of each vertex u satisfies $(\boldsymbol{x}(u), \boldsymbol{y}(u), \boldsymbol{z}(u)) \leq (n, 2n, 2n)$.*

Next we define the slope of an arc based only in the coordinates of its incident vertices. Let $\mathbf{sx}(e)$ be the slope of line segment joining u and v after it has been projected onto the plane $x = 0$; and similarly for $\mathbf{sy}(e)$ and $\mathbf{sz}(e)$. It is not hard to ensure that the coordinates of any vertex are nonzero, however to use these slope functions as defined, we do not want any arcs parallel to any axis. Hence, we requre one extra property of the graph embedding to ensure that these slopes are always bounded and nonzero. If $e = uv$ is an edge of G, then $\mathbf{x}(u) \neq \mathbf{x}(v)$, $\mathbf{y}(u) \neq \mathbf{y}(v)$, and $\mathbf{z}(u) \neq \mathbf{z}(v)$. We conjecture that the embedding given in [8] can be modified in one of two ways to accomodate this, if so we call this a modified grid embedding. Either we allow rational vertex coordinates instead of integers, or we increase the given bound of $(n, 2n, 2n)$ by a constant amount. We propose the following identifier for each arc $e = (u, v)$ of $\vec{G}$.

$$
\begin{aligned}
\mathbf{key}(e) = \quad & (\mathbf{x}(\mathbf{smaller}(e)), \mathbf{y}(\mathbf{smaller}(e)), \mathbf{z}(\mathbf{smaller}(e)), \\
& \mathbf{sx}(e), \mathbf{sy}(e), \mathbf{sz}(e), \\
& \mathbf{x}(u), \mathbf{y}(u), \mathbf{z}(u)).
\end{aligned}
$$

Finally, we define $e < e'$ whenever $\mathbf{key}(e) < \mathbf{key}(e')$ lexographically. The next lemma proves that this defines a total order on the arcs of $\vec{G}$. We conjecture that this also defines a compatible total order.

Lemma 3.4.4. *If G has a modified grid embedding, then the $\mathbf{key}()$ function induces a total order on the arcs of $\vec{G}$.*

Proof. Clearly, the key exists for any arc, and for any pair of arcs, their keys are comparable. Hence, it is sufficient to show distinct arcs have distinct keys. Let

$e = (u, v)$ and $e' = (u', v')$ be distinct arcs of $\vec{G}$. Clearly, if $u \neq u'$, then we can see by the last three coordinates that **key**$(e) \neq$ **key**(e'), (regardless of whether or not e and e' have another vertex in common). From now on suppose $u = u'$. Since $e \neq e'$, this implies $v \neq v'$.

If **smaller**$(e) = v$, then, regardless of whether or not **smaller**(e') is u' or v', we are done, since $v \neq u'$ and $v \neq v'$, and either one implies that the first three coordinates of **key**(e) and **key**(e') are different. Similarly if **smaller**$(e') = v'$ we are done.

Finally suppose **smaller**$(e) = u =$ **smaller**(e'). So the first three, and the last three coordinates of **key**(e) and **key**(e') are the same. Suppose to the contrary that the arcs have equal slope in all three planes, (so the middle three coordinates are equal). So u, v, and v' are colinear, with u on the line segment joining v and v'. Now, since the slopes are always nonzero by supposition, this implies that either **smaller**$(e) \neq u$ or **smaller**$(e') \neq u$. This is a contradiction.

We conclude that **key**() induces a total order on the arcs of $\vec{G}$.

$\square$

Future work would be to prove the aforementioned graph embedding exists, and show this does imply a compatible total order on the arcs of $\vec{G}$.

Bibliography

[1] R. Armoni, A. Ta-Shma, A. Wigderson, S. Zhou. An $O(\log^{4/3} n)$ space algorithm for (s,t) connectivity in undirected graphs. *Journal of the ACM*, **47** 294–311

[2] G. Barnes, W. Ruzzo. Undirected *st*-connectivity in polynomial time and sublinear space. *Comput. Complexity*, **6**(1) 1–28 (1996)

[3] P. Beame, A. Borodin, P. Raghavan, W. Ruzzo, M. Tompa. A time-space tradeoff for undirected graph traversal for walking automata. *SIAM J. Comput.*, **28**(3) 1051–1072 (1999)

[4] P. Bose, P. Morin. An improved algorithm for subdivision traversal without extra storage. *Internat. J. Comput. Geom. Appl.*, **12**(4) 297–308 (2002)

[5] A. Broder, A. Karlin, P. Raghavan, E. Upfal. Trading space for time in undirected *s-t* connectivity. *SIAM J. Comput.*, **23** 324–334 (1994)

[6] E. Chávez, S. Dobrev, E. Kranakis, J. Opatrny, L. Stacho, J. Urrutia. Traversal of a quasi-planar subdivision without using mark bits. *Journal of Interconnection Networks*, **5**(4) (2004)

[7] S. Cook, C. Rackoff. Space lower bounds for maze threadability on restricted machines. *SIAM J. Computing*, **9**(3) 636–652 (1980)

[8] R. Cohen, P. Eades, Tao Lin, F. Ruskey. Three-Dimensional Graph Drawing. *Algorithmica*, **17** 199–208 (1997)

[9] R. Cohen, P. Fraigniaud, D. Ilcinkas, A. Korman, D. Peleg. Label-guided graph exploration by a finite automaton. *ACM Transactions on Algorithms*, **4**(4) 331–344 (2008)

[10] M. de Berg, M. van Kreveld, R. van Oostrum, M. Overmars. Simple traversal of a subdivision without extra storage. *International Journal of Geogrpahic Information Systems*, **11** 349–373 (1997).

[11] S. Dobrev, J. Jansson, K. Sadakane, W.K. Sung. Finding short right-hand-on-the-wall walks in graphs. *Proc. 12th Colloquium on Structural Information and Communication Complexity (SIROCCO)*, **3499** 127–139 (2005)

[12] G. Dudek, M. Jenkin, W. Milios, D. Wilkes. Robotic exploration as graph construction. *IEEE Trans. Robot. Automat.*, **7** 859–865 (1991)

[13] J. Edmonds, Time-space trade-offs for undirected ST-connectivity on a JAG, *Proceesings of the 25th Annual ACM Symposium on Theory of Computing*, 718–727 (1993)

[14] U. Fiege. A spectrum of time-space tradeoffs for undirected *s-t* connectivity. *J. Comput. System Sci.*, **54** 305–316 (1997)

[15] R. Fleischer. The square of every two-connected graph is Hamiltonian. *Journal of Combinatorial Theory, Series B*, **16**(1) 29–34 (1974)

[16] P. Fraigniaud, D. Ilcinkas, G. Peer, A. Pelc, D. Peleg. Graph exploration by a finite automaton. *Theoretical Computer Science*, **345** 331–344 (2005)

[17] L. Gąsieniec, R. Klasing, R. Martin, A. Navarra, X. Zhang. Fast periodic graph exploration with constant memory. *J Computer and System Science*, **74**(5) 808–822 (2008)

[18] A. Georgakopoulos. A short proof of Fleischner's theorem. *Discrete Mathematics*, **309** 6632–6634 (2009)

[19] F. Jaeger. A survey of the cycle double cover conjecture. *Annals of Discrete Mathematics*, **27** 1–12 (1985)

[20] F. Jaeger. Flows and generalized coloring theorems in graphs. *J. Combinatorial Theory Ser. B*, **26** 205–216 (1979)

[21] D. Ilcinkas. Setting port numbers for fast graph exploration. *Theoretical Computer Science*, **401** 236–242 (2008)

[22] E. Nissan, E. Szemeredi, A. Wigderson. Undirected connectivity in $O(\log^{1.5} n)$ space. *Proc. of 33rd Annual Symposium on Foundations of Computer Science, IEEE*, 24–49 (1992)

[23] H. Rollik, Automaten in planaren graphen. *Acta Informatica*, **13** 287–298 (1980)

[24] P. Seymour. Disjoint paths in graphs. *Discrete Mathematics*, **29**(3) 293–309 (1980)

[25] G. Szekeres. Polyhedral decomposition of cubic graphs. *Bulletin of the Australian Mathematical Society*, **8** 367–387 (1973)